

Working with 

A University of Queensland Advanced Workshop

Session 2: Objects and Their Manipulation

Bill Venables, CSIRO/Data61, Dutton Park

Rhetta Chappell, Griffith University

2–5 February, 2021

Contents

1	Objects and their attributes	3
1.1	Atomic and recursive objects	4
2	A real pain in the neck	13
2.1	Reading in the data	14
2.2	Key functions	16
2.3	Missing values – to plug or not to plug?	18
2.4	Wide form to long form	20
2.5	Looking at the data	23
3	Another example: How fast do you speak?	33
3.1	A modelling investigation	41
3.2	Diagnostics	44

3.3 The robust alternative	47
--------------------------------------	----

Session information	49
----------------------------	-----------

1 Objects and their attributes

"In **S** *everything* is an object and *every* object has a *class*."
John Chambers, c. 1996.

R is a *language for manipulating objects*.
(me)

1.1 Atomic and recursive objects

There are five kinds of *atomic* object. In a sense, all other objects are made from these.

numeric Vectors of numerical quantities.

character Vectors of *character strings*.

logical Vectors with components either *TRUE* or *FALSE*.

complex Vectors of complex numbers (each component a pair of numerical quantities).

raw Vectors of raw bites, mainly used with compiled code.

All have the extra possibility of a *missing value marker* in place of a component. This is shown either as *NA* or *<NA>*.

Other types of objects are known as *recursive*

list objects: *data.frame*, fitted model objects, loads of other stuff....

date objects and their allies: *POSIXct*, *POSIXlt*, *Date*,

language objects: *formula*, *call*, *expression*,

environment objects: the most mysterious of all!

Even *atomic* objects are more than they seem.

- All objects have two *intrinsic attributes*, namely *mode* and *length*.
- Objects can have many *assigned* attributes as well, such as *class*, *names*, *names*, *dimnames*, &c.
- Attributes *self-describe* the object.
- Object attributes can *allow* some operations and *inhibit* others.

Examples

```
set.seed(12345)
x <- runif(16) %>% round(2)
dim(x) <- c(4,4)
dimnames(x) <- list(rows = letters[1:4], columns = LETTERS[1:4]); x[1:2, ]
```

	columns			
rows	A	B	C	D
a	0.72	0.46	0.73	0.74
b	0.88	0.17	0.99	0.00

```
attributes(x)
```

```
$dim
```

```
[1] 4 4
```

```
$dimnames
```

```
$dimnames$rows
```

```
[1] "a" "b" "c" "d"
```

```
$dimnames$columns
```

```
[1] "A" "B" "C" "D"
```

```
c(length(x), mode(x), class(x))
```

```
[1] "16"      "numeric" "matrix"  "array"
```

```
tx <- as.table(x) %>% as.data.frame(responseName = "values")
```

```
tx          ## turn a numeric matrix into a 'long form' data frame
```

	rows	columns	values
1	a	A	0.72
2	b	A	0.88
3	c	A	0.76
....			
14	b	D	0.00
15	c	D	0.39
16	d	D	0.46

```
class(tx)
```

```
[1] "data.frame"
```

```
x <- as.vector(x) ## clear all applied attributes
```

```
attributes(x)
```

```
NULL
```



```

f <- factor(fill0(round(10*x)))  ## see note to follow on fill0(...)
table(f)

f
00 02 03 04 05 07 08 09 10
 2  2  1  1  3  3  1  2  1

attributes(f); mode(f)

$levels
[1] "00" "02" "03" "04" "05" "07" "08" "09" "10"

$class
[1] "factor"
[1] "numeric"

f+1      ## arithmetic with factors is inhibited!

[1] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA

```

Turning a factor back into a numeric vector: two possibilities

```
as.character(f)
```

```
[1] "07" "09" "08" "09" "05" "02" "03" "05" "07" "10" "00" "02" "07" "00"  
[15] "04" "05"
```

```
rbind(A = as.numeric(f),  
      B = as.numeric(as.character(f))) %>%  
  booktabs(digits=0)
```

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
A	6	8	7	8	5	2	3	5	6	9	1	2	6	1	4	5
B	7	9	8	9	5	2	3	5	7	10	0	2	7	0	4	5

Filling with zeros

- Filling up numerical strings with initial zeros to a constant number of digits can preserve lexicographic ordering in line with numeric.
- The [WWRCourse](#) package has two functions that behave identically, but are coded very differently:

```
##  
## Method 1: using string substitution  
fill0 <- function(x) {  
  gsub(" ", "0", format(x, justify = "right"))  
}  
##  
## Method 2: using character size computations and paste  
zfill <- function(x) {  
  m <- max(n <- nchar(x <- as.character(x)))  
  paste0(strrep(0, m-n), x)  
}
```

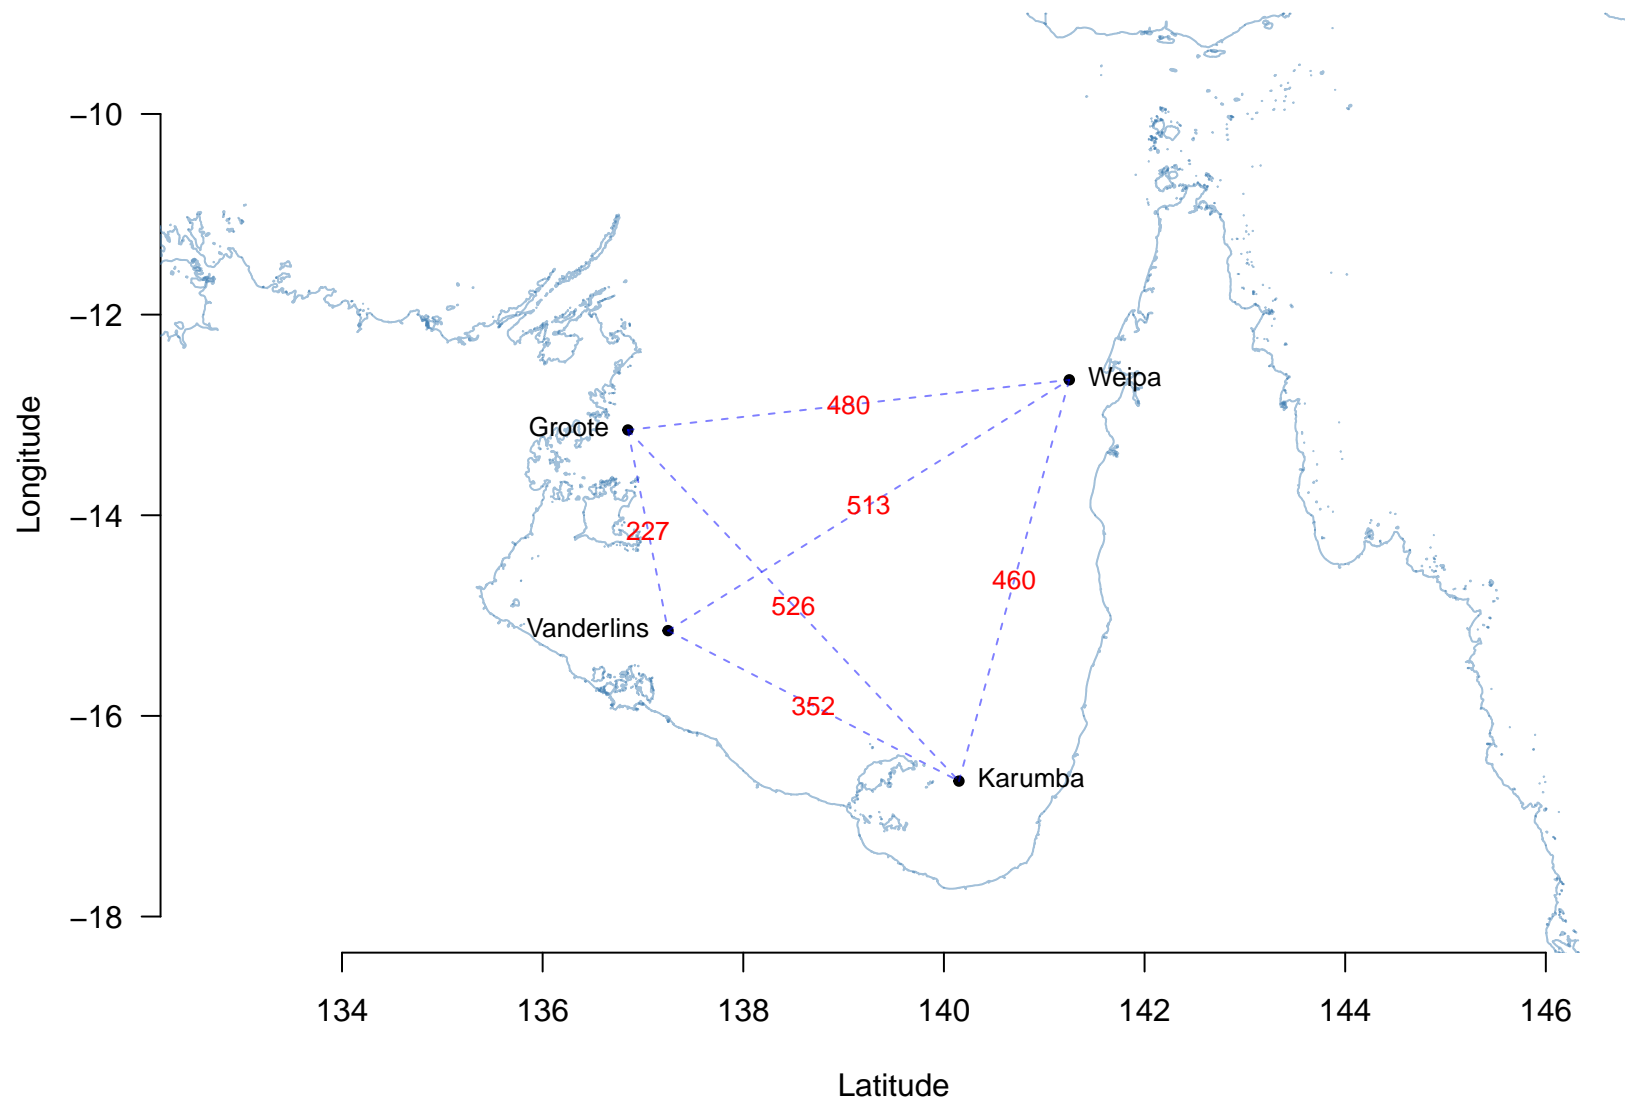
- Make sure you know how both work!

Complex types

The forgotten atomic type?

For four given locations in the Gulf of Carpentaria, find the shortest marine distances between each of them and show on a map.

```
with(data4, {  
  blueish <- alpha("steelblue", 0.5)  
  z <- complex(real = Longitude, imaginary = Latitude)  
  plot(z, asp = 1, xlim = c(135, 144), ylim = c(-18, -9),  
       bty = "n", xlab = "Latitude", ylab = "Longitude")  
  lines(0z, col = blueish)  
  text(z, Place, pos = c(2,2,4,4), cex=0.8)  
  ij <- utils::combn(4,2)  
  i <- ij[1, ]  
  j <- ij[2, ]  
  segments(z[i], z[j], lty = "dashed", col = alpha("blue", 0.5))  
  text((z[i]+z[j])/2, round(gcd_km(z[i], z[j])), col="red", cex=0.8)  
})
```



2 A real pain in the neck

- We use some “typical” data^a on the outcome of a physiotherapy project to study two ways of ameliorating neck pain, mainly in office workers.
- Two groups:
 - Control: Given advice on neck pain management
 - Intervention: Given an exercise programme and advice.
- Five assessment times: 0, 3, 6, 9 and 12 months. A *longitudinal* study.
- Several explanatory variables: Age, Sex, Occupation, &c.
- Data supplied as a [stata](#) binary, available in the [WWRCourse](#) package as extra data folder, *extdata/stata*.

^aKindly supplied by Dr Xiaoqi Chen, UQ

2.1 Reading in the data

- The modern tool collection, (“*tidyverse*”), supplies a number of packages for specific data input jobs:

readr For reading mainly text files, especially .csv files.

readxl For reading mainly **Microsoft** Excel files.

haven For reading a range of foreign system files, *including stata binaries*.

- All three produce as the result of their main functions a peculiarly enhanced *data.frame* known (annoyingly) as a “tibble”^a.
Fortunately it is easy to turn a tibble back into a normal `data.frame`, and for the most part, we shall.
- We will use the “chaining” operator, *%>%*, (often incorrectly called a ‘pipe’, but we do it too!) throughout this session without further comment, for clarity.

^aNo, it’s not a cat.

```

library(haven)
NPStudy <- read_stata(system.file("extdata", "stata", "NeckPainStudy.dta",
                                package = "WWRCourse"))

sapply(NPStudy, class)

$idid
[1] "numeric"

$grp
[1] "haven_labelled" "vctrs_vctr"      "double"
....
$pain_num_9m
[1] "numeric"

$pain_cat_9m
[1] "numeric"

```

labelled columns will become character;
 numeric columns need to be cleared.

2.2 Key functions

The *dplyr* package supplies five *key* functions, namely

filter For selecting rows (cf. *subset*)

select For selecting columns (cf. *subset(..., select = ...)*)

mutate (sic!) For computations on columns (cf. *within()*)

arrange For ordering rows of a data set

summarise Together with *group_by*, for summarising data sets in various ways. (cf. *apply()*.)

do While not a key function, often useful. (cf, *by()* in old speak)

```

NPStudy <- read_stata(system.file("extdata", "stata", "NeckPainStudy.dta",
                                package = "WWRCourse")) %>%

  lapply(function(x) if(is.labelled(x)) {
    as_factor(x)
  } else {
    as_vector(x)
  }) %>% untibble()
Store(NPStudy)  ## for safe keeping

```

```

NeckPain <- NPStudy %>%
  select(-contains("_cat_")) %>% ## remove categoricals
  within({
    Ident <- paste0("S", fill0(idd)) %>% factor()
    Cluster <- paste0("C", fill0(cluster)) %>% factor()
  }) %>%
  rename(Sex = sex, Group = grp, Organisation = organisation,
         Industry = industry, Ergo = ergo, Age = age,
         BMI = bmi, Education = edu, Occupation = occ,
         Comorbidity = cm) %>%
  select(Ident, Cluster, Group, Organisation:Occupation,
         Comorbidity, pain_num_b:pain_num_9m)

```

2.3 Missing values – to plug or not to plug?

```
colSums(is.na(NeckPain))
```

Ident	Cluster	Group	Organisation	Industry
0	0	0	0	0
Ergo	Age	BMI	Sex	Education
30	6	2	0	0
Occupation	Comorbidity	pain_num_b	pain_num_3m	pain_num_12m
0	0	0	150	357
pain_num_6m	pain_num_9m			
394	466			

These are a problem. In an exploratory analysis it seems useful:

- To plug missing values, if there are not too many, in the *explanatory* variables *only*.
- To do this, use *only* information in the *explanatory* variables, **not** the responses.
- Don't forget this is *only* for exploratory purposes!

We will use an imputation process based on random forest technology.
This is provided in the workshop software.

```
set.seed(20200211)
Extract <- NeckPain %>% select(Group:Comorbidity) %>% ## only these
  rfImputeUnsupervised() %>% ## The black box plugger
  within({ ## do a bit of cleaning-up
    Ergo <- round(Ergo) ## make things look innocent
    Age <- round(Age)
    BMI <- round(BMI, 2)
  })
NeckPain[, names(Extract)] <- Extract ## plug the gaps
rm(Extract) ## destroy the evidence...
colSums(is.na(NeckPain)) ## check all is OK
```

Ident	Cluster	Group	Organisation	Industry
0	0	0	0	0
Ergo	Age	BMI	Sex	Education
0	0	0	0	0
Occupation	Comorbidity	pain_num_b	pain_num_3m	pain_num_12m
0	0	0	150	357
pain_num_6m	pain_num_9m			
394	466			

2.4 Wide form to long form

The *tidyr* package is used for re-shaping data sets, most commonly “wide form” to or from “long form”.

Two main functions: *pivot_longer* and *pivot_wider*.

```
longNeckPain <- NeckPain %>%
  pivot_longer(pain_num_b:pain_num_9m,
               names_to = "Time", values_to = "NPain") %>%
  within({
    time <- Time %>% ## currently a mess
      sub("b$", "0", .) %>% ## final 'b' -> 0
      gsub("[^[:digit:]]", "", .) %>% ## ditch any non-digit
      as.numeric() ## coerce to a number
    Time <- ordered(paste0("T", fill0(time)))
    Treat <- ifelse(time == 0, "Base", paste0(substring(Group, 0, 1),
                                              substring(Time, 2)))
  }) %>% select(Ident:Group, Time, time, Treat,
               Organisation:Comorbidity, NPain) %>%
  arrange(Ident, Time) %>% na.omit() %>% untibble()
Store(NeckPain, longNeckPain)
```

What have we got? Let's see a few bits:

```
longNeckPain %>%
```

```
  select(Ident, Group:Treat, Ergo:Sex, NPain)
```

	Ident	Group	Time	time	Treat	Ergo	Age	BMI	Sex	NPain
1	S001	Control	T00	0	Base	27	50	26.56	female	2.0
2	S001	Control	T03	3	C03	27	50	26.56	female	0.0
3	S002	Intervention	T00	0	Base	28	33	23.15	female	0.0
4	S002	Intervention	T03	3	I03	28	33	23.15	female	0.0
5	S002	Intervention	T06	6	I06	28	33	23.15	female	0.0
....										
2328	S761	Control	T06	6	C06	37	30	29.40	female	2.0
2329	S761	Control	T09	9	C09	37	30	29.40	female	6.0
2330	S761	Control	T12	12	C12	37	30	29.40	female	1.0
2331	S762	Control	T00	0	Base	31	26	21.60	male	7.0
2332	S763	Control	T00	0	Base	34	38	34.05	female	5.0
2333	S763	Control	T03	3	C03	34	38	34.05	female	3.0

What is the *Treat* factor?

```
with(longNeckPain,  
      tapply(Treat, list(Group, Time), unique)) %>%  
      booktabs()                                ## for a neat result!
```

	T00	T03	T06	T09	T12
Control	Base	C03	C06	C09	C12
Intervention	Base	I03	I06	I09	I12

It reflects the fact that at time 0 no intervention has been applied at all: both groups supply a baseline neck pain level to which the other “treatments” may be compared.

2.5 Looking at the data

Ignore most explanatory variables and look at numbers and means, first.


```
meanNPain <- longNeckPain %>%  
  group_by(Group, Time, Treat, Sex) %>%  
  summarise(N = n(), Pain = mean(NPain), .groups = "drop") %>%  
  ungroup() %>%  
  untibble() %>%  
  arrange(Treat, Sex)
```

meanNPain

	Group	Time	Treat	Sex	N	Pain
1	Control	T00	Base	male	143	1.4755245
2	Intervention	T00	Base	male	146	0.9383562
3	Control	T00	Base	female	230	1.9086957
4	Intervention	T00	Base	female	221	1.8506787
....						
17	Intervention	T09	I09	male	44	1.1363636
18	Intervention	T09	I09	female	69	1.3768116
19	Intervention	T12	I12	male	66	0.7424242
20	Intervention	T12	I12	female	111	1.7297297


The mean neck pain score is low, (fortunately!), but how does it vary?

```
sex_cols <- c(male = "steelblue", female = "rosybrown")
ggplot(meanNPain) +
  aes(x = Time, y = Pain, colour = Sex, group = Sex) +
  geom_point(aes(size = N)) +
  geom_line() +
  facet_wrap(~Group, ncol=2) +
  theme(legend.position = "bottom") +
  ylim(0, 2) +
  scale_colour_manual(values = sex_cols)
```



To put it the other way:

```
ggplot(meanNPain) + aes(x = Time, y = Pain, colour = Sex,  
                        linetype = Group, group = Group) +  
  geom_point(aes(size = N)) + geom_line() + facet_wrap(~Sex, ncol=2) +  
  theme(legend.position = "bottom") + ylim(0, 2) +  
  scale_colour_manual(values = sex_cols)
```




How many people stuck with the programme?

```
meanNPain %>%  
  pivot_wider(id_cols = c(Group, Sex), ## These columns identify rows  
              names_from = Time,  
              values_from = N) %>% booktabs
```

Group	Sex	T00	T03	T06	T09	T12
Control	male	143	115	68	55	78
Intervention	male	146	113	63	44	66
Control	female	230	192	122	106	128
Intervention	female	221	170	93	69	111

Some drop-off mid-programme, but some return at the end of the year. A picture makes this much easier to notice:

```
mNPain <- meanNPain %>% within({  
  Time <- factor(as.character(Time), levels = rev(levels(Time)))  
})  
ggplot(mNPain) + aes(x = Time, y = N, fill = Sex) +  
  geom_bar(stat = "identity", position = "dodge") + coord_flip() +  
  facet_wrap(~ Group) + theme(legend.position = "bottom") +  
  scale_fill_manual(values = sex_cols)
```




Rather than mean pain level, the researchers were more interested in whether or not participants exceeded a pain threshold score of 3 or more.

```
percentNPain <- longNeckPain %>%  
  group_by(Group, Time, Treat, Sex) %>%  
  summarise(N = n(), Severe = mean(NPain >= 3) * 100,  
            .groups = "drop") %>%  
  ungroup() %>%  
  arrange(Treat, Sex) %>%  
  untibble()  
percentNPain %>%  
  pivot_wider(id_cols = Group:Treat, names_from = Sex,  
              values_from = Severe) %>% booktabs()
```


Group	Time	Treat	male	female
Control	T00	Base	23.08	33.04
Intervention	T00	Base	13.01	32.58
Control	T03	C03	21.74	25.00
Control	T06	C06	16.18	22.13
Control	T09	C09	18.18	25.47
Control	T12	C12	23.08	27.34
Intervention	T03	I03	18.58	20.59
Intervention	T06	I06	6.35	18.28
Intervention	T09	I09	15.91	21.74
Intervention	T12	I12	7.58	28.83

```
ggplot(percentNPain) + aes(x = Time, y = Severe, colour = Sex, group = Sex) +
  geom_point(aes(size = N)) + geom_line() + facet_wrap(~Group, ncol=2) +
  theme(legend.position="bottom") + ylab("% score 3 or higher") +
  ylim(0, 40) + scale_colour_manual(values = c(male = "steelblue",
                                                female = "rosybrown"))
```



Or, to put it another way:

```
ggplot(percentNPain) + aes(x = Time, y = Severe, colour = Sex,  
                           linetype = Group, group = Group) +  
  geom_point(aes(size = N)) + geom_line() + facet_wrap(~Sex, ncol=2) +  
  theme(legend.position="bottom") + ylab("% score 3 or higher") +  
  ylim(0, 40) + scale_colour_manual(values = sex_cols)
```



3 Another example: How fast do you speak?

We use a data set described and provided in the package [hqmisc](#).^a Our version is called *talkers*.

From the help information on the package:

talkers	package:hqmisc	R Documentation
A data frame with 80 observations on the following 6 variables.		
'id' identifier code (from data source, see Source)		
'sex' sex (0=female, 1=male)		
'age' age (in years)		
'region' region of origin (a factor with levels 'M'=Mid, 'N'=North, 'S'=South, or 'W'=West)		
'syldur' average duration of syllables, or seconds per syllable (in seconds, excluding pause time, 1/(articulation rate))		
'nsyl' average number of syllables per phrase, or average phrase length in syllables		

^aWhich is not included in the set of packages needed for this course.

We begin by setting up a more convenient version of the data.

```
dutchSpeakers <- talkers %>% within({
  id <- paste0("S", zfill(id))    ## function in WWRSoftware
  sex <- recode_factor(sex, `0` = "Female", `1` = "Male")
  AgeGroup <- factor(ifelse(age > 42, "45+", "40-"),
                     levels = c("40-", "45+"))
  region <- factor(as.character(region),
                  levels = c("N",      "M",      "W",      "S"),
                  labels = c("North", "Middle", "West", "South"))
}) %>%
  rename(Ident = id, Region = region, Sex = sex, ## New = old
         Draw1 = syldur, Verbose = nsyl) %>%
  arrange(Region, Sex, age)
```


It is important to check this all went according to plan.

```
dutchSpeakers   ### look at it
```

	Ident	Sex	age	Region	Drawl	Verbose	AgeGroup
1	S119	Female	28	North	0.2856	7.85	40-
2	S124	Female	28	North	0.2500	12.47	40-
3	S125	Female	32	North	0.3132	7.63	40-
4	S134	Female	33	North	0.2543	6.96	40-
5	S123	Female	36	North	0.2650	10.79	40-
6	S120	Female	47	North	0.3109	7.55	45+
....							
75	S114	Male	40	South	0.2293	12.35	40-
76	S148	Male	48	South	0.2357	6.17	45+
77	S147	Male	51	South	0.2904	6.13	45+
78	S144	Male	56	South	0.2330	6.19	45+
79	S149	Male	57	South	0.2741	6.71	45+
80	S142	Male	59	South	0.2927	5.45	45+


Look at the data graphically. How *Drawl* vary with *Sex*, *Region* and *AgeGroup*?

```
ggplot(dutchSpeakers) + aes(x = age, y = Drawl, colour = Sex) +  
  geom_point() +  
  geom_smooth(method = "lm", formula = y ~ (x < 43)) +  
  theme(legend.position = "top") +  
  geom_vline(xintercept = c(41, 45), linetype = "longdash",  
            colour = "dark green") + xlab("\nAge of respondent") +  
  facet_wrap( ~ Region) +  
  scale_colour_viridis_d()
```



Does the drawl depend on the prolixity?

```
p0 <- ggplot(dutchSpeakers) + aes(Verbose, Drawl) +  
  geom_point(colour="brown") +  
  geom_smooth(method="lm", formula=y~x, size=0.7)  
p0 + facet_grid(Sex ~ Region)
```




Generally, drawl falls with prolixity: people using longer phrases tend to utter their individual syllables more quickly.

There is a very slow-speaking female from the South region.

Further break-up of the sample by age group will result in groups of size 5, which may be rather small to see patterns.

Or will it?

```
ggplot(dutchSpeakers) + aes(x = Verbose, y = Draw1) +  
  geom_point(colour = "brown", size = 1) +  
  geom_smooth(method = "lm", se = FALSE, fullrange = TRUE, size=0.5,  
              formula = y~x) +  
  facet_grid(AgeGroup + Sex ~ Region)
```

Young people talk faster than old?

Males speak more quickly than females?

Regional differences are present, but hard to categorize?

3.1 A modelling investigation

```
dsModel <- aov(Drawl ~ Verbose+Sex+AgeGroup+Region, dutchSpeakers)
anova(dsModel) %>%
  booktabs(digits = c(0,0,4,4,2,4)) ## for display only. Cut these bits
```

Df	Sum Sq	Mean Sq	F value	Pr(>F)
1	0.0144	0.0144	18.57	0.0001
1	0.0095	0.0095	12.32	0.0008
1	0.0066	0.0066	8.54	0.0046
3	0.0139	0.0046	6.01	0.0010
73	0.0564	0.0008		

Because the design is very close to orthogonal, the order in which the terms enter the model is not important.

```
dropterm(dsModel)  ## WWRUtilities enhanced version
```

Single term deletions

Model:

```
Draw1 ~ Verbose + Sex + AgeGroup + Region
```

	Df	Sum of Sq	RSS	AIC	F Value	Pr(F)
<none>			0.056402	-566.58		
Verbose	1	0.0059815	0.062384	-560.52	7.7418	0.0068635
AgeGroup	1	0.0067715	0.063174	-559.51	8.7642	0.0041414
Sex	1	0.0102496	0.066652	-555.22	13.2658	0.0005023
Region	3	0.0139359	0.070338	-554.92	6.0123	0.0010189

Are the coefficients consistent with what we saw in the graphics?

```
round(summary.lm(dsModel)$coeff, 4) ## check on signs and sizes
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.3061	0.0146	20.8989	0.0000
Verbose	-0.0041	0.0015	-2.7824	0.0069
SexMale	-0.0231	0.0063	-3.6422	0.0005
AgeGroup45+	0.0186	0.0063	2.9604	0.0041
RegionMiddle	-0.0180	0.0088	-2.0408	0.0449
RegionWest	-0.0336	0.0088	-3.8188	0.0003
RegionSouth	-0.0038	0.0089	-0.4271	0.6705





```
#### manual residual analysis
```

```
diagDutchSpeakers <- data.frame(rs = scale(resid(dsModel)),  
                                 fv = fitted(dsModel))
```

3.2 Diagnostics


For diagnostic checks, traditional graphics are simpler and usually adequate. The simplest way is to “plot” the fitted model object:

```
layout(rbind(1:2, 3:4))  ## 2 x 2 array of plots, filled by rows  
plot(dsModel)
```



We see what's going on a bit better if we do it by hand:

```
layout(rbind(1:2))
par(bg = "#F4FCFE", family="mono")  ## mundane, old-fashioned look...
plot(rs ~ fv, diagDutchSpeakers, col = "steel blue",
     xlab = "fitted values", ylab = "scaled residuals",
     main = "Residuals vs Fitted Values")
abline(h = 0, col = "red", lty = "dashed")
with(diagDutchSpeakers, {
  qqnorm(rs, xlab = "normal scores", col = "steel blue",
        ylab = "sorted scaled residuals",
        main = "Normal Scores Plot")
  qqline(rs, col = "red", lty = "dashed")
})
```



The slow speaking person shows up as a clear outlier.
We should check the extent to which our results are influenced by it.
We could omit it and see if anything changes.
A better approach is to use a robust method and see what it suggests.

3.3 The robust alternative

With statistical modelling:

- You usually believe the **data**, and the challenge is to construct a **model** which will separate the true signals from the noise as well as possible.
- With robust methods, you usually have a good idea of what an appropriate **model** should be, but you have grave doubts about some of the **data**. The challenge is to identify the *consistent core of good data* so that your model estimates, and hence your conclusions, are not misleading.
- Using both methods in tandem involves a philosophical inconsistency and needs much caution!
- There is no fully reliable way to identify outliers from the data itself alone: *context is always important*.

Session information

Date: 2021-01-29

- R version 4.0.3 (2020-10-10), x86_64-pc-linux-gnu
- Running under: Ubuntu 20.04.1 LTS
- Matrix products: default
- BLAS: /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.9.0
- LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.9.0
- Base packages: base, datasets, graphics, grDevices, methods, stats, utils
- Other packages: dplyr 1.0.3, english 1.2-5, forcats 0.5.1, ggplot2 3.3.3, ggthemes 4.2.4, gridExtra 2.3, haven 2.3.1, knitr 1.31, lattice 0.20-41, patchwork 1.1.1, purrr 0.3.4, readr 1.4.0, scales 1.1.1, stringr 1.4.0, tibble 3.0.5, tidyr 1.1.2, tidyverse 1.3.0, WWRCourse 0.2.3, WWRData 0.1.0, WWRGraphics 0.1.2, WWRUtilities 0.1.2, xtable 1.8-4
- Loaded via a namespace (and not attached): assertthat 0.2.1, backports 1.2.1, broom 0.7.3, cellranger 1.1.0, cli 2.2.0, colorspace 2.0-0, compiler 4.0.3, crayon 1.3.4, DBI 1.1.1, dbplyr 2.0.0, digest 0.6.27, ellipsis 0.3.1, evaluate 0.14, fansi 0.4.2, farver 2.0.3, fractional 0.1.3, fs 1.5.0, generics 0.1.0, glue 1.4.2, grid 4.0.3, gtable 0.3.0, highr 0.8, hms 1.0.0, http 1.4.2, iterators 1.0.13, jsonlite 1.7.2, labeling 0.4.2, lazyData 1.1.0, lifecycle 0.2.0, lubridate 1.7.9.2, magrittr 2.0.1, MASS 7.3-53, Matrix 1.3-2, mgcv 1.8-33, modelr 0.1.8, munsell 0.5.0, nlme 3.1-151, parallel 4.0.3, PBSmapping 2.73.0, pillar 1.4.7, pkgconfig 2.0.3, R6 2.5.0, randomForest 4.6-14, Rcpp 1.0.6, readxl 1.3.1, reprex 1.0.0, rlang 0.4.10, rpart 4.1-15, rstudioapi 0.13, rvest 0.3.6, SOAR 0.99-11, splines 4.0.3, stringi 1.5.3, tidyselect 1.1.0, tools 4.0.3, vctrs 0.3.6, viridisLite 0.3.0, withr 2.4.1, xfun 0.20, xml2 1.3.2