

Working with 

A University of Queensland Advanced Workshop

Session 5: Elementary Linear Models

Bill Venables, CSIRO/Data61, Dutton Park

Rhetta Chappell, Griffith University

2–5 February, 2021

Contents

1	A simple example: the Janka Hardness data	2
1.1	Alternative parametrizations	5
1.2	Diagnostics and corrective action	9
1.3	A transformation?	10
1.4	A simpler approach	19
1.5	Bootstrap confidence intervals for the mean	24
1.5.1	Classical version	24
1.5.2	Bayesian version	27
1.6	Some more recent data	30
	References	34
	Session information	35

1 A simple example: the Janka Hardness data

Example taken from E. J. Williams (1959) *Regression Analysis*.

Two variables:

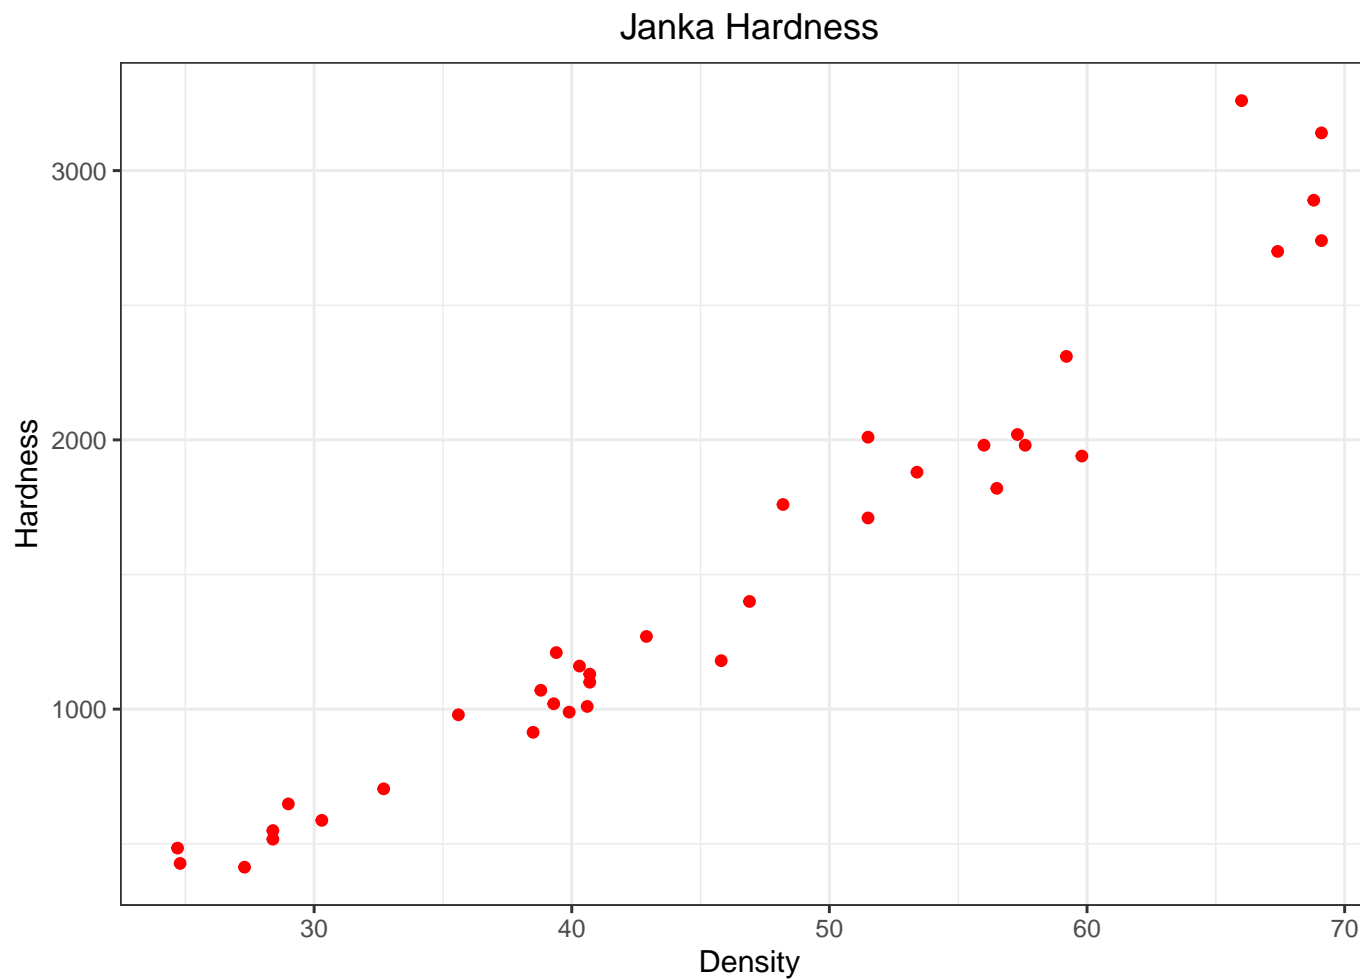
Hardness : The Janka hardness of a sample of timbers, (in lbs). (See [here](#) for a discussion.)

Density : The density of the sample, (in lbs/ft^3)

The problem: *Build a predictor of Hardness from Density.*

```
theme_set(theme_bw() + theme(plot.title = element_text(hjust=0.5)))
```

```
ggplot(janka) + aes(x = Density, y = Hardness) +  
  geom_point(colour = "red") + labs(title = "Janka Hardness")
```



- Clearly strong dependence of Hardness on Density;
- Possibly curvilinear – try polynomials first;
- Possibly with unequal variance?

```
m1 <- lm(Hardness ~ Density, janka)
m2 <- update(m1, . ~ . + I(Density^2))
m3 <- update(m2, . ~ . + I(Density^3))
round(summary(m3)$coef, 4)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-641.4379	1235.6550	-0.5191	0.6073
Density	46.8637	86.3018	0.5430	0.5909
I(Density^2)	-0.3312	1.9140	-0.1730	0.8637
I(Density^3)	0.0060	0.0135	0.4405	0.6625

1.1 Alternative parametrizations

Nothing is “significant”! Why?

A modification:

```
m1a <- lm(Hardness ~ I(Density - 50), janka)
m2a <- update(m1a, . ~ . + I((Density - 50)^2))
m3a <- update(m2a, . ~ . + I((Density - 50)^3))
round(summary(m3a)$coef, 4)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1618.6529	43.4597	37.2449	0.0000
I(Density - 50)	58.4367	4.8619	12.0193	0.0000
I((Density - 50)^2)	0.5626	0.1999	2.8147	0.0083
I((Density - 50)^3)	0.0060	0.0135	0.4405	0.6625

The models are fully equivalent, but the coefficients refer to different things

The explanation in graphical terms

```
m0a <- lm(Hardness ~ 1, janka)  ## constant predictor model
m4a <- update(m3a, .~.+I((Density-50)^4))

pJanka <- data.frame(Density = -5:75)
pJanka <- within(pJanka, {
  constant <- predict(m0a, pJanka)
  linear <- predict(m1a, pJanka)
  quadratic <- predict(m2a, pJanka)
  cubic <- predict(m3a, pJanka)
  quartic <- predict(m4a, pJanka)
})

pJankaLong <- pJanka %>%
  pivot_longer(names_to="Model", values_to="Hardness", constant:quartic) %>%
  mutate(Model = factor(Model, levels = c(constant, linear, quadratic,
                                           cubic, quartic)))

titlex <- expression("Model:"*' $H = \beta_0 + \beta_1(D - \alpha) + \beta_2(D - \alpha)^2 + \cdots + \epsilon$ '*italic(H) == ' $\beta_0 + \beta_1(D - \alpha) + \beta_2(D - \alpha)^2 + \cdots + \epsilon$ '*italic(beta[0] +
  beta[1]*(D - alpha) + beta[2]*(D - alpha)^2 + cdots + epsilon))
```


```

p <- ggplot(janka) + aes(x = Density, y = Hardness) + geom_point() +
  geom_line(aes(colour = Model), data = pJankaLong, size = 0.5) +
  geom_hline(yintercept = 0, linetype = "solid", size = 0.25) +
  geom_vline(xintercept = c(0,50), linetype = "solid",
             size = 0.25, colour = "blue") +
  annotate("text", x = c(0,50)-1, y = 3500, hjust = 1,
          label = paste("alpha ==", c(0,50)), parse = TRUE,
          size = 4) +
  labs(x = expression(italic(D):' Density'),
       y = expression(italic(H):' Hardness'),
       title = titlex) +
  theme(text = element_text(size = 12), legend.position = "bottom",
        plot.title = element_text(hjust = 0.5, vjust = 1, size = 12)) +
  scale_colour_brewer(palette = "Set1")

```





p

$$\text{Model: } H = \beta_0 + \beta_1(D - \alpha) + \beta_2(D - \alpha)^2 + \dots + \varepsilon$$



1.2 Diagnostics and corrective action

```
layout(matrix(1:4, 2, 2, byrow = TRUE))  
plot(m2a)
```



1.3 A transformation?


The main purpose of a transformation is to provide a scale in which the response is homoscedastic — but not necessarily the only purpose.

Transforming the response will affect both the mean structure and the variance.


Consider a Box-Cox transformation on a) the straight line model and b) the quadratic model.

```
layout(matrix(1:2, 1))  
box_cox(m1a)  
title(main = "straight line model", line = 3)  
box_cox(m2a)  
title(main = "quadratic model", line = 3)
```

straight line model



quadratic model





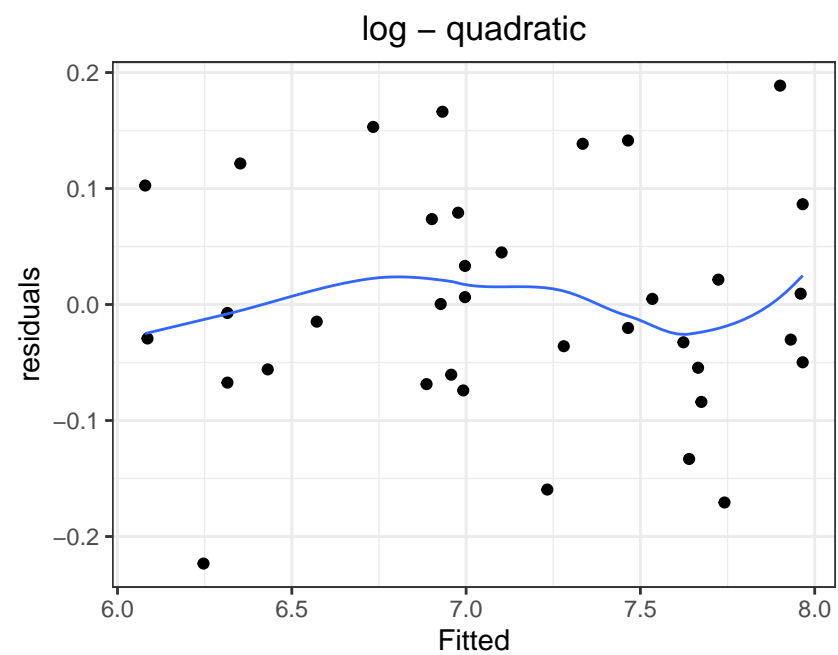
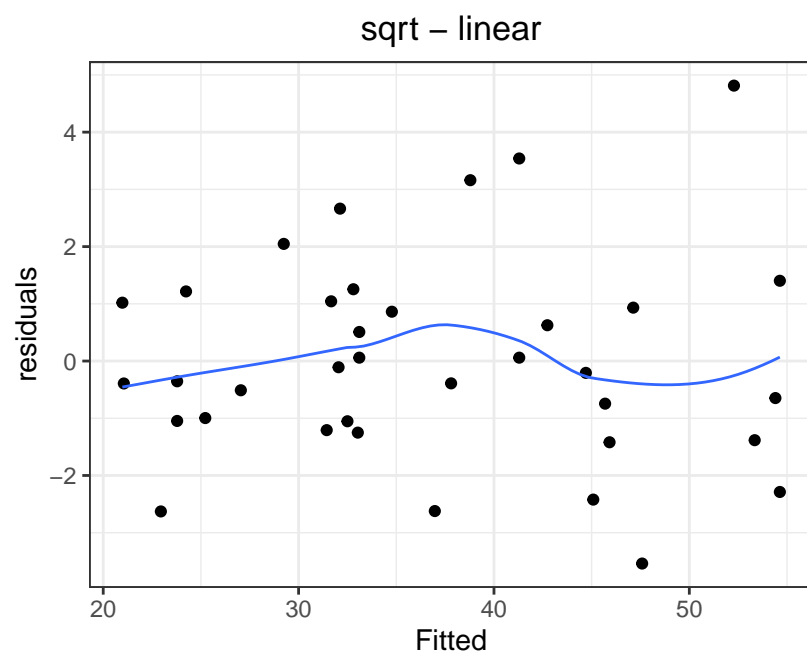
Assuming a straight line — a square root transformation?

Allowing for a quadratic response — a log transformation?

Consider the effect of the transformation on residuals:

```
m1 <- lm(sqrt(Hardness) ~ poly(Density, 1), janka)
m2 <- lm(log(Hardness) ~ poly(Density, 2), janka)
Janka <- within(janka, {
  r1 <- resid(m1);   r2 <- resid(m2)
  f1 <- fitted(m1);  f2 <- fitted(m2)
})
p0 <- ggplot(Janka)
p1 <- p0 + aes(x = f1, y = r1) + geom_point() +
  labs(x = "Fitted", y = "residuals", title = "sqrt - linear") +
  geom_smooth(se = FALSE, method = "loess", size = 0.5, formula = y~x)
p2 <- p0 + aes(x = f2, y = r2) + geom_point() +
  labs(x = "Fitted", y = "residuals", title = "log - quadratic") +
  geom_smooth(se = FALSE, method = "loess", size = 0.5, formula = y~x)
p3 <- p0 + aes(sample = r1) + stat_qq() + stat_qq_line()
p4 <- p0 + aes(sample = r2) + stat_qq() + stat_qq_line()

(p1 + p2)/(p3 + p4)
```



The message so far:

- A square root transformation straightens out the regression line, but
- A log transformation is necessary to even out the variance.

This suggests a generalized linear model:

- with a *sqrt* link, to give a scale in which the response is straight line,
- with a variance function $\propto \mu^2$ to allow for variance heterogeneity.

Within the GLM family this suggests a Gamma model. The *sqrt* link is non-standard, which *used* to be a problem, but no longer.


```

mGLM <- glm(Hardness ~ I(Density-50),
            family = Gamma(link = "sqrt"),
            data = janka)
mGLM2 <- update(mGLM, .~.+I((Density-50)^2))
round(summary(mGLM2)$coef, 4)

```


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	40.4191	0.4382	92.2296	0.0000
I(Density - 50)	0.7516	0.0313	24.0035	0.0000
I((Density - 50)^2)	-0.0013	0.0017	-0.7373	0.4661

The straight line model seems adequate – residual checks:


```
rs <- resid(mGLM)
fv <- fitted(mGLM)
layout(matrix(1:2, 1))
plot(fv, rs, xlab = "fitted values", ylab = "residuals",
     pch = 20, col=grey(0.5), main = "Variance uniformity")
abline(h = 0, lty = "dashed", col="red"); grid()

qqnorm(rs, pch = 20, col = grey(0.5), xlab = "Normal scores",
       ylab = "sorted residuals", main = "Normal Q-Q plot")
qqline(rs, col="red", lty = "dashed"); grid()
```

Variance uniformity



Normal Q-Q plot




1.4 A simpler approach

- What happens if we transform *both* response and predictor?
- A multiplicative model seems heuristically reasonable:

$$H = \alpha D^\beta \exp E \quad \implies \quad \log H = \alpha^* + \beta \log D + E \quad (\alpha^* = \log \alpha)$$

```
ggplot(janka) + aes(x = Density, y = Hardness) +  
  geom_point() + scale_x_log10() + scale_y_log10() +  
  stat_smooth(method = "lm", size = 0.5, colour = "hot pink", formula = y~x)
```

- This suggests a Gamma model with a *log* link and a log-transformed predictor, linear term only. So it proves to be.
- Using a GLM, leaving the response untransformed, allows us to predict on the natural scale, thus avoiding the complications associated with back transforming predictions.



Compare the final model with a simple quadratic regression, ignoring variance heterogeneity.

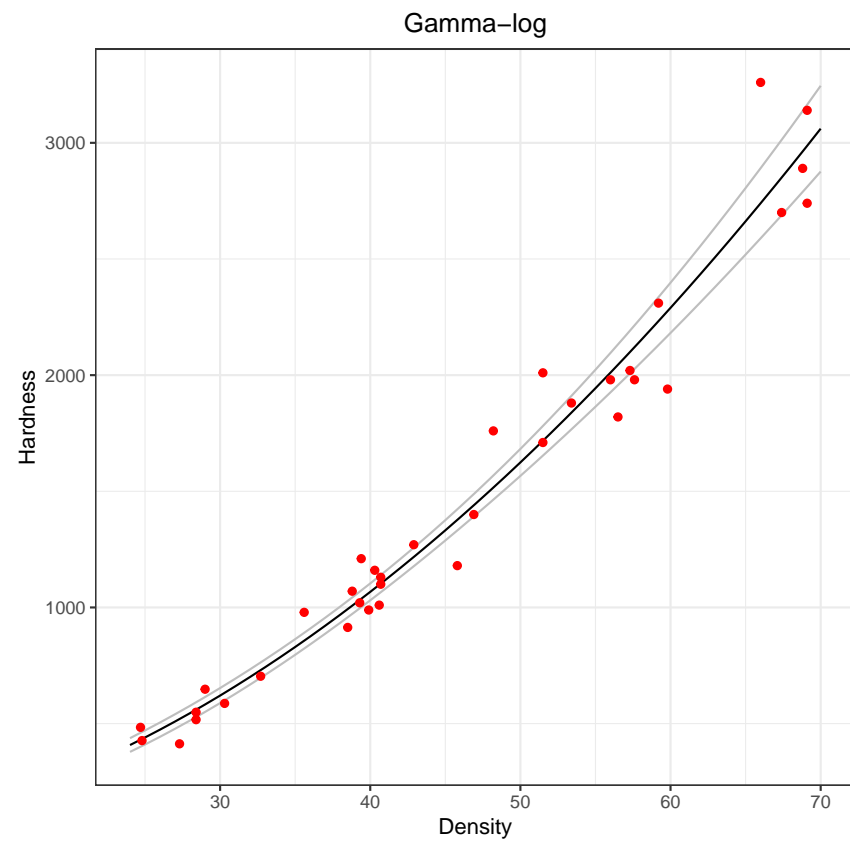
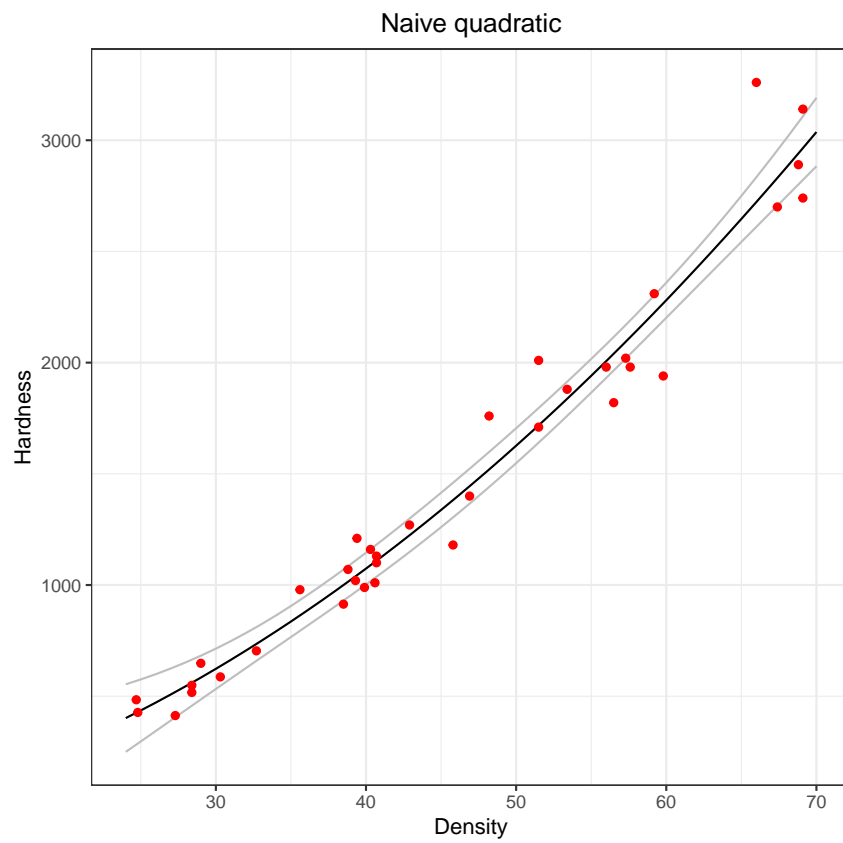
```
jankaLM <- lm(Hardness ~ poly(Density, 2), janka)
pJanka <- data.frame(Density = 24:70)
ext <- predict(jankaLM, pJanka, type = "resp", se.fit = TRUE) %>%
  as.data.frame() %>%
  within({
    lower <- fit - 2*se.fit
    upper <- fit + 2*se.fit
  })
pJanka <- cbind(pJanka, ext)
pLM <- ggplot(pJanka) + aes(x = Density) +
  geom_line(aes(y = fit), colour = "black") +
  geom_line(aes(y = lower), colour = "grey") +
  geom_line(aes(y = upper), colour = "grey") +
  geom_point(data = janka, aes(y = Hardness), colour = "red") +
  xlab("Density") + ylab("Hardness") + labs(title = "Naive quadratic")
```

```

jankaGLM <- glm(Hardness ~ log(Density), Gamma(link = log), janka)
pJanka <- data.frame(Density = 24:70)
ext <- predict(jankaGLM, pJanka, type = "resp", se.fit = TRUE) %>%
  as.data.frame() %>%
  within({
    lower <- fit - 2*se.fit
    upper <- fit + 2*se.fit
  })
pJanka <- cbind(pJanka, ext)

pGLM <- ggplot(pJanka) + aes(x = Density) +
  geom_line(aes(y = fit), colour = "black") +
  geom_line(aes(y = lower), colour = "grey") +
  geom_line(aes(y = upper), colour = "grey") +
  geom_point(data = janka, aes(y = Hardness), colour = "red") +
  xlab("Density") + ylab("Hardness") + labs(title = "Gamma-log")
#
# gridExtra::grid.arrange(pLM, pGLM, nrow = 1)
pLM + pGLM

```



1.5 Bootstrap confidence intervals for the mean

1.5.1 Classical version

The idea is that we predict the mean for a large number of *bootstrap samples* of the original model, re-fitting the model as if the bootstrap sample were the real data. The quantiles of the bootstrap predictions provide the appropriate confidence interval.

The code is simple and explains the method more precisely.

```

set.seed(20210202)
boot_sample <- function(data) data[sample(nrow(data), replace = TRUE), ]


ci <- replicate(500, {
  tmp <- update(jankaGLM, data = boot_sample(janka)) ## bootstrap data
  predict(tmp, pJanka, type = "resp") ## predictions for target
}) %>%
  apply(1, quantile, prob = c(0.025, 0.975))

pJanka <- pJanka %>%
  within({
    lowerBB <- ci[1, ]
    upperBB <- ci[2, ]
  })

greenish <- alpha("dark green", 0.5)
pGLM + labs(title = "Classical bootstrap") +
  geom_line(data = pJanka, aes(x = Density, y = lowerBB),
            colour = greenish, size = 0.5) +
  geom_line(data = pJanka, aes(x = Density, y = upperBB),
            colour = greenish, size = 0.5)

```

Classical bootstrap



1.5.2 Bayesian version

An alternative to the classical bootstrap is the Bayesian Bootstrap idea of Rubin (1981).

- Re-fit the model with *random weights* for the observations.
- Choosing $W \sim \text{Exp}(1)$ gives $E W = 1 = \text{Var } W$, the same as for the normal bootstrap. (Rubin gives a theoretical justification.)

```
X <- matrix(sample(100, size = 1000*100, replace = TRUE), ncol = 100) %>%  
  apply(1, tabulate, nbins = 100)  
CB <- c(mean      = mean(colMeans(X)),  
        variance = mean(colMeans((X-1)^2)))  
W <- rexp(10000)  
BB <- c(mean = mean(W), variance = var(W))  
rbind(Classical = CB, Bayesian = BB)
```

	mean	variance
Classical	1.0000000	0.991460
Bayesian	0.9958835	1.005921

The code is nearly identical to the classical case.

```

set.seed(20210202)
Norm <- function(x) x/mean(x)


ci <- replicate(500, {
  tmp <- update(jankaGLM, weights = Norm(rexp(nrow(janka))))
  predict(tmp, pJanka, type = "resp")
}) %>%
  apply(1, quantile, prob = c(0.025, 0.975))

pJanka <- pJanka %>%
  within({
    lowerBB <- ci[1, ]
    upperBB <- ci[2, ]
  })

greenish <- alpha("dark green", 0.5)
pGLM + labs(title = "Bayesian bootstrap") +
  geom_line(data = pJanka, aes(x = Density, y = lowerBB),
            colour = greenish, size = 0.5) +
  geom_line(data = pJanka, aes(x = Density, y = upperBB),
            colour = greenish, size = 0.5)

```


Bayesian bootstrap




1.6 Some more recent data

These date from 2012. They were obtained from the internet, but the units were not given.

```
Janka <- within(Janka2012, {  
  Type <- ifelse(grepl("Eucalyptus", Binomial), "Eucalpyt", "Other")  
})  
  
jp <- ggplot(Janka) + aes(x = Density, y = Hardness) + geom_point() +  
  aes(colour = Type) + # coord_trans("log10", "log10") +  
  scale_x_log10() + scale_y_log10() +  
  scale_colour_brewer(palette = "Set1", name = "Genus")  
jp  
jp + ## identify the light species  
  ggrepel::geom_text_repel(data = filter(Janka, Density < 600),  
    aes(label = Name), size = 3,  
    colour = greenish, ## some transparency  
    fontface = "bold", family = "serif") +  
  geom_smooth(method = "lm", se = FALSE, formula = y ~ x, size = 0.5)
```





Comparing model estimates:

```
oldJanka <- glm(Hardness~log(Density), Gamma(link="log"), janka)
newJanka <- update(oldJanka, data = Janka)
format(cbind(old = coef(oldJanka),
              recent = coef(newJanka)),
       digits = 5) %>% booktabs()
```

	old	recent
(Intercept)	0.026034	-10.605491
log(Density)	1.883123	1.893879

The discrepancy in the intercept coefficients is due to a mismatch of units.
The coefficient of log(Density) is unit free.

References

Rubin, D. B. (1981). The bayesian bootstrap. *The Annals of Statistics* 9, 130–134.

Venables, W. N. and B. D. Ripley (2002). *Modern Applied Statistics with S* (Fourth ed.). New York: Springer. ISBN 0-387-95457-0.

Session information

Date: 2021-01-29

- R version 4.0.3 (2020-10-10), x86_64-pc-linux-gnu
- Running under: Ubuntu 20.04.1 LTS
- Matrix products: default
- BLAS: /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.9.0
- LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.9.0
- Base packages: base, datasets, graphics, grDevices, methods, stats, utils
- Other packages: dplyr 1.0.3, english 1.2-5, forcats 0.5.1, ggplot2 3.3.3, ggthemes 4.2.4, gridExtra 2.3, knitr 1.31, lattice 0.20-41, patchwork 1.1.1, purrr 0.3.4, readr 1.4.0, scales 1.1.1, stringr 1.4.0, tibble 3.0.5, tidyr 1.1.2, tidyverse 1.3.0, WWRCourse 0.2.3, WWRData 0.1.0, WWRGraphics 0.1.2, WWRUtilities 0.1.2, xtable 1.8-4
- Loaded via a namespace (and not attached): assertthat 0.2.1, backports 1.2.1, broom 0.7.3, cellranger 1.1.0, cli 2.2.0, colorspace 2.0-0, compiler 4.0.3, crayon 1.3.4, DBI 1.1.1, dbplyr 2.0.0, digest 0.6.27, ellipsis 0.3.1, evaluate 0.14, fansi 0.4.2, farver 2.0.3, fractional 0.1.3, fs 1.5.0, generics 0.1.0, ggrepel 0.9.1, glue 1.4.2, grid 4.0.3, gtable 0.3.0, haven 2.3.1, highr 0.8, hms 1.0.0, httr 1.4.2, iterators 1.0.13, jsonlite 1.7.2, labeling 0.4.2, lazyData 1.1.0, lifecycle 0.2.0, lubridate 1.7.9.2, magrittr 2.0.1, MASS 7.3-53, Matrix 1.3-2, mgcv 1.8-33, modelr 0.1.8, munsell 0.5.0, nlme 3.1-151, parallel 4.0.3, PBSmapping 2.73.0, pillar 1.4.7, pkgconfig 2.0.3, R6 2.5.0, randomForest 4.6-14, RColorBrewer 1.1-2, Rcpp 1.0.6, readxl 1.3.1, reprex 1.0.0, rlang 0.4.10, rpart 4.1-15, rstudioapi 0.13, rvest 0.3.6, SOAR 0.99-11, splines 4.0.3, stringi 1.5.3, tidyselect 1.1.0, tools 4.0.3, vctrs 0.3.6, withr 2.4.1, xfun 0.20, xml2 1.3.2