

Working with 

A University of Queensland Advanced Workshop

Session 7: Random Effects Extensions

Bill Venables, CSIRO/Data61, Dutton Park

Rhetta Chappell, Griffith University

2–5 February, 2021

Contents

1	An introductory example: petroleum extraction	2
1.1	Fixed or random?	6
2	An extended example: going fishing	15
2.1	A brief look at generalized linear/additive mixed models	24
3	Count response variables	34
3.1	Negative Binomial models: the quine example	34
3.2	Gladstone Bream recruitment index	40
A	Theme functions	49
B	Two helper functions	49
	References	52
	Session information	53

1 An introductory example: petroleum extraction

The petrol data of N. L. Prater.

- *No* crude oil sample identification label. (Factor.)
- *SG* specific gravity, degrees API. (Constant within sample.)
- *VP* vapour pressure in pounds per square inch. (Constant within sample.)
- *V10* volatility of crude; ASTM 10% point. (Constant within sample.)
- *EP* desired volatility of gasoline. (The end point. Varies within sample.)
- *Y* yield as a percentage of crude.

For a description in **R**:

```
?petrol  
petrol  ## print the whole data!
```

	No	SG	VP	V10	EP	Y
1	A	50.8	8.6	190	205	12.2
2	A	50.8	8.6	190	275	22.3
3	A	50.8	8.6	190	345	34.7
4	A	50.8	8.6	190	407	45.7
....						
28	I	32.2	2.4	284	351	14.0
29	I	32.2	2.4	284	424	23.2
30	J	31.8	0.2	316	365	8.5
31	J	31.8	0.2	316	379	14.7
32	J	31.8	0.2	316	428	18.0

For a more complete description of the data and an alternative (somewhat fussy) analysis see the **betareg** package, (Cribari-Neto and Zeileis, 2010). *?GasolineYield*.

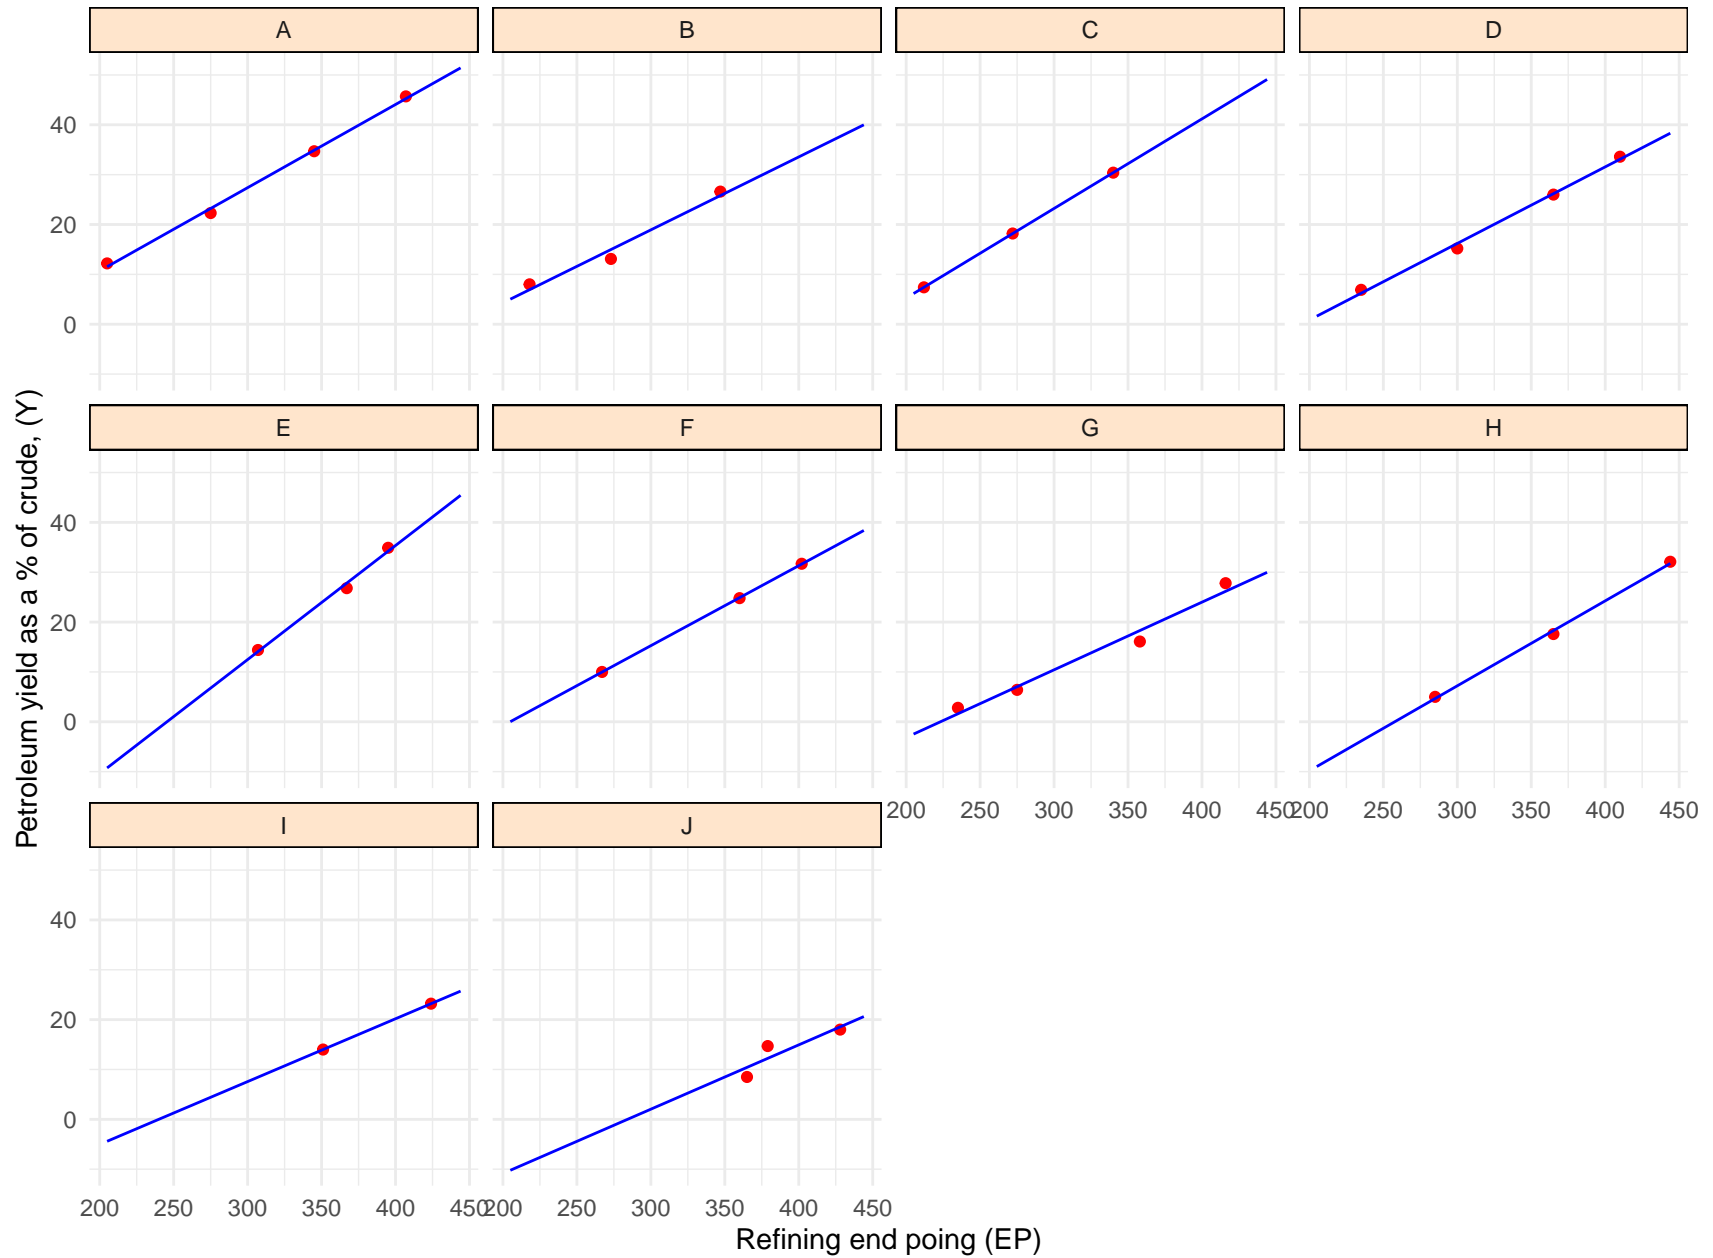
An initial look at the data:

```
ggplot(petrol) + aes(x = EP, y=Y) +  
  geom_point(colour = "red") +  
  stat_smooth(method = "lm", se = FALSE, fullrange = TRUE,  
              size=0.5, colour = "blue", formula=y~x) + facet_wrap( ~ No) +  
  labs(title = "Petrol refining data of N. L. Prater",  
        x = "Refining end poing (EP)",  
        y = "Petroleum yield as a % of crude, (Y)")
```

The plots are shown below It will be convenient later to have a centred version of *EP*:

```
petrol <- petrol %>%  
  within(EPc <- scale(EP)) ### for convenience  
Store(petrol)
```

Petrol refining data of N. L. Prater



1.1 Fixed or random?

A pure fixed effects model treats the crude oil samples as independent with the residual error as the only source of randomness.

A random effects model treats them as possibly dependent, in that they may share the value of a latent random variable, addition to the residual error.

The obvious candidate predictor to be regarded as injecting an additional source of randomness is the crude oil sample indicator, *No*.

Fixed effects only.

```
options(show.signif.stars = FALSE)
m3 <- lm(Y ~ 0 + No/EPc, petrol)          ## 10 ints + 10 slopes
m2 <- lm(Y ~ 0 + No+EPc, petrol)         ## 10 ints + 1 slope
m1 <- lm(Y ~ 1 + SG+VP+V10+EPc, petrol)  ## (1 int + 3 coeffs) + 1 slope
anova(m1, m2, m3)
```

Analysis of Variance Table

Model 1: Y ~ 1 + SG + VP + V10 + EPc

Model 2: Y ~ 0 + No + EPc

Model 3: Y ~ 0 + No/EPc

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	27	134.804				
2	21	74.132	6	60.672	4.0009	0.01965
3	12	30.329	9	43.803	1.9257	0.14390

Parallel regressions, but differences between samples cannot quite be explained by regression on the other variables.

Random effects alternatives:

```
requireData(lme4)      ## alt. nlme
Rm1 <- lmer(Y ~ 1 + SG+VP+V10 + EPc + (1|No),      data = petrol)
Rm2 <- lmer(Y ~ 1 + SG+VP+V10 + EPc + (1+EPc|No), data = petrol)
anova(Rm1, Rm2)        ## automatic re-fitting
```

refitting model(s) with ML (instead of REML)

Data: petrol

Models:

Rm1: Y ~ 1 + SG + VP + V10 + EPc + (1 | No)

Rm2: Y ~ 1 + SG + VP + V10 + EPc + (1 + EPc | No)

	npar	AIC	BIC	logLik	deviance	Chisq	Df	Pr(>Chisq)
Rm1	7	149.38	159.64	-67.692	135.38			
Rm2	9	153.34	166.54	-67.672	135.34	0.0385	2	0.9809

Emphatically different slopes are not needed!

NB:

- Default fitting method “*REML*” produces good estimates of variance components
- To compare models using LR, models must be fitted with method “*ML*”
- *anova* detects if this has happened and updates the object if necessary.

Inspecting the random effects fit:

```
print(summary(Rm1), correlation = FALSE)
```

```
Linear mixed model fit by REML ['lmerMod']
```

```
Formula: Y ~ 1 + SG + VP + V10 + EPc + (1 | No)
```

```
Data: petrol
```

```
REML criterion at convergence: 143.9
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-1.7807	-0.6064	-0.1068	0.4572	1.7811

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
No	(Intercept)	2.087	1.445
Residual		3.505	1.872

```
Number of obs: 32, groups: No, 10
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	46.06264	14.69241	3.135

SG	0.21940	0.14694	1.493
VP	0.54586	0.52053	1.049
V10	-0.15424	0.03996	-3.860
EPc	10.96402	0.38980	28.128

```
print(summary(Rm2), correlation = FALSE)
```

Linear mixed model fit by REML ['lmerMod']

Formula: Y ~ 1 + SG + VP + V10 + EPc + (1 + EPc | No)

Data: petrol

REML criterion at convergence: 143.9

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.77775	-0.59744	-0.07372	0.46067	1.76997

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
No	(Intercept)	2.13847	1.4624	
	EPc	0.06278	0.2506	0.06
Residual		3.43304	1.8528	

Number of obs: 32, groups: No, 10

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	45.93487	14.77738	3.108
SG	0.21923	0.14770	1.484

VP	0.55213	0.52349	1.055
V10	-0.15380	0.04023	-3.823
EPc	10.96966	0.39677	27.648

Use *fixef* for fixed effect estimates and *ranef* for BLUPs:

```
cbind(Rm1 = ranef(Rm1)$No, Rm2 = ranef(Rm2)$No)
```

	(Intercept)	Rm2.(Intercept)	Rm2.EPc
A	-0.05943867	-0.04207529	0.05301543
B	-0.21857133	-0.21900738	-0.02073023
C	1.92029758	1.96452890	0.01941871
D	-1.92760817	-1.95646887	-0.03046897
E	-0.21650182	-0.22737149	0.06972143
F	0.56931165	0.57476410	0.01571855
G	0.06701596	0.05019416	-0.10419793
H	0.19194473	0.18788703	0.04550772
I	-0.40278068	-0.40719019	-0.03461858
J	0.07633074	0.07473900	-0.01336613

Variances and correlations

```
VarCorr(Rm2)          ## an unhelpful print method.
```

Groups	Name	Std.Dev.	Corr
No	(Intercept)	1.46235	
	EPc	0.25056	0.057
Residual		1.85285	

```
names(VarCorr(Rm2)) %>% noquote()

[1] No

VarCorr(Rm2)[["No"]] ## The pieces are all accessible
```

	(Intercept)	EPc
(Intercept)	2.13847440	0.02098379
EPc	0.02098379	0.06278021

```
attr(,"stddev")
```

	(Intercept)	EPc
	1.4623524	0.2505598

```
attr(,"correlation")
```

	(Intercept)	EPc
(Intercept)	1.00000000	0.05726913
EPc	0.05726913	1.00000000

2 An extended example: going fishing

The *Headrope* data set gives catch and effort data from a prawn fishery.

- The fishery has 7 *Stock* regions *Tig1*, ..., *Tig7*, West to East.
- The data is for 20 seasons (*YearF*) 1987, ..., 2006. (*Y2K* = year - 2000.)
- There are 236 *Vessels*, which visit one or more stock regions within a season, each for one or more *Days*.
- The *response* for which a model is required is the total *Catch* in kg, by a vessel within a stock region for a season.
- Additionally the vessels have *Hull* size, engine *Power* and the *Headrope* length they were using recorded. (These are constant within season, but may change between seasons.)


```
dim(Headrope)

[1] 8594    13

head(Headrope, 2)

      YearF Y2K Stock Vessel Days Head Hull Power Catch Banana Tiger
0001  1987 -13  Tig1   V008   20   20  133   350  4355   2509   975
0002  1987 -13  Tig1   V012   13   20  134   336  4746   3612   252
      Endeavour King
0001           871    0
0002           882    0

Headrope <- Headrope %>% within(YearF <- factor(YearF)) ## needed
Store(Headrope)
```

The purpose of the study was to gain some insight on the marginal effect of headrope length on the catch.

A multiplicative (log-linear) model was suggested, with additive random effects for a) vessel and b) stock regions over seasons.

Two random effects models: the first is the simpler

```
HRmodel1 <- lmer(log(Catch) ~ 0 + log(Days) + Y2K + log(Head) +
  log(Power) + log(Hull) + Stock +
  (1|Vessel) + (1|YearF/Stock), data = Headrope,
  control = lmerControl(optimizer = "bobyqa"))
HRmodel1_ML <- update(HRmodel1, REML = FALSE)
Store(HRmodel1, HRmodel1_ML, lib = .Robjjects)
```

The second has a more elaborate random effect structure:

```
HRmodel2 <- lmer(log(Catch) ~ 0 + log(Days) + Y2K + log(Head) +
  log(Power) + log(Hull) + Stock +
  (1|Vessel) + (0+Stock|YearF), data = Headrope,
  control = lmerControl(optimizer = "bobyqa"))

boundary (singular) fit: see ?isSingular

HRmodel2_ML <- update(HRmodel2, REML = FALSE)

boundary (singular) fit: see ?isSingular

Store(HRmodel2, HRmodel2_ML, lib = .Robjjects)
```

The more elaborate model seems justified by AIC, but not BIC!

```
anova(HRmodel11_ML, HRmodel12_ML)
```

Data: Headrope

Models:

HRmodel11_ML: $\log(\text{Catch}) \sim 0 + \log(\text{Days}) + \text{Y2K} + \log(\text{Head}) + \log(\text{Power}) + \log(\text{Hull})$

HRmodel11_ML: $\text{Stock} + (1 \mid \text{Vessel}) + (1 \mid \text{YearF}/\text{Stock})$

HRmodel12_ML: $\log(\text{Catch}) \sim 0 + \log(\text{Days}) + \text{Y2K} + \log(\text{Head}) + \log(\text{Power}) + \log(\text{Hull})$

HRmodel12_ML: $\text{Stock} + (1 \mid \text{Vessel}) + (0 + \text{Stock} \mid \text{YearF})$

	npar	AIC	BIC	logLik	deviance	Chisq	Df	Pr(>Chisq)
HRmodel11_ML	16	11676	11788	-5821.8	11644			
HRmodel12_ML	42	11613	11910	-5764.5	11529	114.48	26	4.501e-13

The fixed effects estimates are very similar:

```
cbind(m1 = fixef(HRmodel1), m2 = fixef(HRmodel2))
```

	m1	m2
log(Days)	1.16175473	1.16092771
Y2K	0.02742599	0.03477483
log(Head)	0.30269966	0.30166004
log(Power)	0.11566567	0.11414072
log(Hull)	0.20684792	0.20812242
StockTig1	2.84888065	2.88900893
StockTig2	2.39961501	2.43790967
StockTig3	2.14663311	2.18115993
StockTig4	2.32495254	2.35986363
StockTig5	2.41038427	2.44534551
StockTig6	2.55091856	2.56126561
StockTig7	2.19453754	2.17294182

Some notes:

- The coefficient on *log(Days)* is slightly larger than 1, (but significantly). A coefficient of 1 would imply that, *mutatis mutandis*, catch is proportional to “effort” (measured in boat days).
- The coefficient of *Y2K* suggests an average fishing power increase in the order of 2.5%–3.5% per year. This looks about right, but it is confounded with change in the stock abundance. Essentially the job of disentangling this confounding is what stock assessment is all about (and why it is so hard).

For reference we include a copy of the summary of the more elaborate model below.

```
print(summary(HRmodel2), correlation = FALSE)
```

Linear mixed model fit by REML ['lmerMod']

Formula:

$$\log(\text{Catch}) \sim 0 + \log(\text{Days}) + \text{Y2K} + \log(\text{Head}) + \log(\text{Power}) + \log(\text{Hull}) + \text{Stock} + (1 \mid \text{Vessel}) + (0 + \text{Stock} \mid \text{YearF})$$

Data: Headrope

Control: lmerControl(optimizer = "bobyqa")

REML criterion at convergence: 11593

Scaled residuals:

Min	1Q	Median	3Q	Max
-16.3368	-0.3270	0.0071	0.4091	4.9784

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
Vessel	(Intercept)	0.01028	0.1014	
YearF	StockTig1	0.16147	0.4018	
	StockTig2	0.11352	0.3369	0.30
	StockTig3	0.02282	0.1511	0.11 0.38
	StockTig4	0.01923	0.1387	0.20 0.76 0.55
	StockTig5	0.02021	0.1422	0.16 0.62 0.62 0.93

StockTig6	0.15939	0.3992	-0.07	0.15	0.36	0.22	0.53
StockTig7	0.17844	0.4224	0.03	0.35	0.26	0.06	0.30
Residual	0.21124	0.4596					

0.80

Number of obs: 8594, groups: Vessel, 236; YearF, 20

Fixed effects:

	Estimate	Std. Error	t value
log(Days)	1.160928	0.004801	241.786
Y2K	0.034775	0.004265	8.153
log(Head)	0.301660	0.049988	6.035
log(Power)	0.114141	0.037573	3.038
log(Hull)	0.208122	0.031699	6.566

```
StockTig1  2.889009    0.195859    14.750
StockTig2  2.437910    0.188492    12.934
StockTig3  2.181160    0.175105    12.456
StockTig4  2.359864    0.174618    13.514
StockTig5  2.445346    0.175125    13.963
StockTig6  2.561266    0.193472    13.238
StockTig7  2.172942    0.196103    11.081
optimizer (bobyqa) convergence code: 0 (OK)
boundary (singular) fit: see ?isSingular
```


2.1 A brief look at generalized linear/additive mixed models

Software for GLMMs is still somewhat developmental.

- *glmmPQL* in *MASS* is based on *nlme*, but handles general cases.
- *glmer* from the *lme4* package handles some GLMMs but is restricted in the families it can take. (In particular, *quasipoisson* is NOT included.)

The software for GAMMs also uses a linear ME engine.

- *gamm* from the *mgcv* package uses *nlme* engine,
- *gamm4* from the *gamm4* package (Wood, 2011), uses *lme4* engine, (and so has the same limitations).

Both *gamm* and *gamm4* return a composite object with an *lme* and a *gam* component. Manipulation is tricky.

To illustrate, we construct a GLMM and a GAMM for the Tiger Prawn species split example. The model structure is slightly simplified relative to the working model.

We use two helper functions, *Hyear* and *twoWay* which will be defined at the end.

First, the GLMM:

```
TModelGLMM <- glmmPQL(Psem/Total ~ ns(Coast, 6) + ns(Sea, 5) +  
  twoWay(DayOfYear, Sea) + ns(Depth, k=5) +  
  Hyear(DayOfYear, 2) + ns(Mud, k=5),  
  random = ~1|Survey,  
  family = quasibinomial, data= Tigers,  
  niter = 40, weights = Total)
```

iteration 1

iteration 2

iteration 3

iteration 4

```
Store(TModelGLMM, lib = .Robjects)
```

Note that the random component is defined separately from the main formula, in [nlme](#) style.

For a GAM with smoothed terms:

```
requireData(mgcv)
TModelGAMM <- gamm(formula = Psem/Total ~ s(Coast, k=5) + s(Sea,k=5) +
                    twoWay(DayOfYear, Sea) +
                    s(DayOfYear, k=5, bs="cc") + s(Depth,k=5) +
                    s(Mud, k=5),
                    knots = list(DayOfYear = seq(0, 364.25, length = 5)),
                    random = list(Survey = ~1),
                    family = quasibinomial, data = Tigers,
                    niterPQL = 40,
                    weights = Total)
```

Maximum number of PQL iterations: 40

iteration 1

iteration 2

iteration 3

iteration 4

iteration 5

iteration 6

iteration 7

```
Store(TModelGAMM, lib = .Robjecs)
```

The random effects from these different models are quite similar. We illustrate below.

```


re1 <- setNames(ranef(TModelGLMM), "GLMM")
re2 <- setNames(ranef(TModelGAMM$lme)$Survey, "GAMM")  ## obscure
(re12 <- cbind(re1, re2))

```

	GLMM	GAMM
Albatross	0.5318898	0.8753641
BSS9708	-1.0162057	-1.4004097
BSS9803	-1.0212852	-1.4685617
CommCatch	0.2927734	0.1931905
CRTryGear	-0.2056465	-0.4231140
DVTryGear	0.5717618	1.0175587
MaximGroote	-0.4686942	-0.6236857
NPFMonitor	0.9254647	0.9378444
Redfield	0.4769554	1.3240899
SpecDist	1.0579519	0.9878992
TEClosure	-0.8708619	-0.9759418
WGoCMonitor	-0.2741036	-0.4442340


Displaying the correspondence “by hand” with base graphics and plot tricks.

```
layout(rbind(1:2), widths = c(3.5,0.5), heights = c(3.5, 3.5),
       respect = TRUE)
with(re12, {
  nos <- seq_along(GLMM)
  lim <- range(GLMM, GAMM)
  plot(GLMM ~ GAMM, pch = 20, col="red", bty="n",
       xlim = lim, ylim = lim, asp = 1)
  abline(0, 1, col = "grey", lty = "dashed")
  box(col = "grey")
  grid()
  pos <- avoid(GAMM, GLMM) ## minimise clashes
  text(GLMM ~ GAMM, labels = nos, xpd = NA, pos=pos)
  par(mar = c(0,0,0,0), xpd = NA)
  frame()
  legend("center", paste(format(nos), rownames(re12),
                        sep = ": "), bty="n")
})
```



But wait, there's more.

```
re12 <- re12 %>% within({
  Survey <- factor(rownames(re1)) %>% reorder(-(GLMM+GAMM), I)
})
ggplot(re12) + aes(x = GAMM, y = GLMM) +
  coord_equal() + xlim(-1.5, 1.2) + ylim(-1.5, 1.2) +
  geom_abline(intercept = 0, slope = 1,
              linetype = "dashed", colour = "grey") +
  geom_point(colour = "red") +
  ggrepel::geom_label_repel(aes(label = Survey),
                           fill = alpha("sky blue", 0.5)) +
  labs(x = "Generalises Additive Mixed Model",
       y = "Generalised Linear Mixed Model") +
  theme_bw() +
  theme(text = element_text(family = "sans"))
```



3 Count response variables

Three most important models for frequency responses are

Binomial: variance less than the mean,

$$V[Y] = \mu(1 - \mu/n).$$

Poisson: variance equal to the mean,

$$V[Y] = \mu.$$

Negative Binomial: variance greater than the mean,

$$V[Y] = \mu(1 + \mu/\theta).$$

It is clear when a Binomial model is appropriate.

Poisson models are mostly used in *surrogate* models for multinomial situations.

3.1 Negative Binomial models: the quine example

Negative Binomial models have proved to be a good option for situations where the response is a count variable and the variance is clearly larger than the mean, but the standard deviation is *roughly proportional* to the mean.

This behaviour often results when the count comes from a process where there are clumps or clusters in the items being counted.

The quine data

- Response: Days away from school by a student in one school year.
- Predictors: Four classifying factors, Age, Lrn, Sex and Eth.

The first clue that this is *not* well modelled as a Poisson response is that the deviance far exceeds the residual degrees of freedom:

```
head(quine, 2)
```

	Eth	Sex	Age	Lrn	Days
1	A	M	FO	SL	2
2	A	M	FO	SL	11

```
q0 <- glm(Days ~ Age*Sex*Eth*Lrn, poisson, quine)
```

```
c(deviance = deviance(q0), "residual d.f." = df.residual(q0))
```


deviance	residual d.f.
1173.899	118.000

Now check the form of any mean-variance relationship:

```
quine %>%
  group_by(Age, Sex, Eth, Lrn) %>%
  summarise(n = n(), mu = mean(Days), sigma = sd(Days), .groups = 'drop') %>%
  na.omit %>%
  unclass %>%
  data.frame -> quine_sum
head(quine_sum, 4)
```

	Age	Sex	Eth	Lrn	n	mu	sigma
1	F0	F	A	AL	4	21.25	17.708284
2	F0	F	N	AL	4	18.50	10.661457
3	F0	M	A	AL	5	13.00	8.031189
4	F0	M	A	SL	3	9.00	6.244998

```
ggplot(quine_sum) + aes(x = mu, y = sigma, size = n) +
  geom_point(colour = "steel blue") +
  geom_smooth(method = "lm", formula = y ~ poly(x, 2), se = FALSE, size = 0.5) +
  geom_smooth(method = "lm", formula = y ~ x, se = FALSE, colour = "red") +
  theme(legend.position = "none")
```



Modelling

Negative binomial model has heuristic and empirical support. Let's fit one and see how it goes.

```
quine_full <- glm.nb(Days ~ Age*Sex*Eth*Lrn, quine)
quine_step <- step_BIC(quine_full) ## BIC penalty
dropterm(quine_step) %>% booktabs(digits = c(0,0,2,2,6))
```

	Df	BIC	LRT	Pr(Chi)
<none>		1132.80		
Age:Sex	3	1136.46	18.62	0.000328
Sex:Eth:Lrn	1	1140.23	12.42	0.000424

(Ex.: check diagnostics to verify things look OK.)

3.2 Gladstone Bream recruitment index

Data from a monitoring survey of bream recruitment in the Gladstone harbour region.

- Four monthly surveys of 26 Sites, December – March
- 20 casts per visit. Response: total bream recruits per visit
- Every site monitored in 2015–16 & 2016–17. Some sporadic visits from 2011–12.
- Some environmental variables recorded, but very roughly.
- Aim: an annual site index for the GHHP health card for each Site.

```
Bream <- subset(GladstoneBream, select = c(Trip, Site:Rock, Casts, Bream))  
names(Bream) %>% noquote
```

```
[1] Trip    Site    Date    Year    Month   Tidal   Depth   Sand    Mud     Gravel  
[11] Rock    Casts   Bream
```

Models

Some experience has led to a set of useful fixed effects and the needs of the index specify a random effect structure. First estimate the NB model including θ .

```
NB_Bream0 <- glmer.nb(Bream ~ offset(log(Casts)) + Depth + Rock + Month +  
                      (1|Site) + (1|Year/Site), data = Bream, nAGQ = 0)  
getME(NB_Bream0, "glmer.nb.theta")  
  
[1] 2.10841
```

There are some advantages in fixing θ at some reasonable value and estimating the parameters of interest with more accuracy.

```
NB_Bream <- glmer(Bream ~ offset(log(Casts)) + Depth + Rock + Month + (1|Site) +  
                  (1|Year/Site), family = negative.binomial(theta = 2),  
                  mustart = fitted(NB_Bream0), data = Bream, nAGQ = 1)
```

Variance components

Our first step is to extract the variance components.

```
(v <- VarCorr(NB_Bream) %>% unlist ) ## these are variances!
```

```
Site:Year      Site      Year  
0.12782639 0.68712447 0.09655783
```

```
sigS <- sqrt(v[["Site"]])  
sigYS <- sqrt(v[["Year"]] + v[["Site:Year"]])  
c(Site = sigS, Year_x_Site = sigYS)
```

```
Site Year_x_Site  
0.8289297 0.4736921
```

The score will be made up of the Year and Site:Year random effects (BLUPs), but we will produce a Site score as well, which gives some idea of the relative productivities of the sites themselves.

```
(RE <- ranef(NB_Bream)) %>% names
```

```
[1] "Site:Year" "Site"      "Year"
```

```
x <- RE[["Site"]]
(S <- data.frame(Site = factor(rownames(x), levels = rownames(x)),
  Score = pnorm(x[["(Intercept)"]], sd = sigS))) %>% head(2)
```

	Site	Score
1	Ramsay Crossing	0.8387646
2	Munduran Creek	0.3076080

```
x <- RE[["Year"]]
(Y <- data.frame(Year = factor(rownames(x), levels = rownames(x)),
  Y_BLUP = x[["(Intercept)"]])) %>% head(2)
```

	Year	Y_BLUP
1	11-12	0.2316966
2	12-13	-0.3604189

```
(x <- RE[["Site:Year"]]) %>% head(3)
```

	(Intercept)
Ramsay Crossing:15-16	0.2985097
Ramsay Crossing:16-17	-0.1459388
Munduran Creek:11-12	0.1878135

This is more involved.

```

YS <- data.frame(nam = rownames(x), stringsAsFactors = FALSE) %>%
  separate(nam, c("Site", "Year"), sep = ":") %>%
  within({
    Site <- factor(Site, levels = levels(S$Site))
    Year <- factor(Year, levels = levels(Y$Year))
    YS_BLUP <- x[["(Intercept)"]]
  }) %>%
  left_join(Y, by = "Year") %>%
  within({
    YS <- pnorm(Y_BLUP + YS_BLUP, sd = sigYS)
    Y_BLUP <- YS_BLUP <- NULL
  }) %>% arrange(Year) %>%
  pivot_wider(names_from = Year, values_from = YS) %>%
  arrange(Site)
booktabs(YS)

```

	Site	11-12	12-13	13-14	14-15	15-16	16-17
1	Ramsay Crossing					0.58	0.62
2	Munduran Creek	0.81	0.24	0.33	0.49	0.37	0.71
3	Black Swan				0.67	0.06	0.97
4	Targinnie Creek	0.67	0.18		0.87	0.16	0.74
5	Graham Creek				0.77	0.31	0.56
6	Hobble Gully				0.56	0.49	0.62
7	Mud Island					0.17	0.79
8	Boat Creek		0.21	0.43	0.80	0.28	0.47
9	Little Enfield Creek				0.76	0.28	0.71
10	Barney Point Pond		0.22	0.44	0.70	0.24	0.57
11	Beecher Creek	0.84	0.12	0.40	0.66	0.29	0.76
12	Old Bruce Highway Bridge				0.45	0.35	0.83
13	Callemondah	0.61	0.08	0.39	0.68	0.49	0.91
14	Farmers Point					0.09	0.91
15	Gatcombe Anchorage					0.34	0.53
16	Wappentake Creek		0.22	0.49	0.59	0.32	0.67
17	South Trees					0.42	0.64
18	Crematorium Pool					0.27	0.85
19	Old Boyne	0.69	0.30		0.61	0.44	0.71
20	Boyne Highway				0.54	0.49	0.78
21	Broadacres					0.36	0.73
22	Iveragh					0.41	0.70
23	Oaky					0.49	0.76
24	7 Mile					0.50	0.75
25	Worthington					0.27	0.69
26	Sandy Bridge					0.39	0.80

Finally, the main result:

```
Scores <- merge(S, YS, by = "Site") %>%  
  within({  
    Site <- factor(as.character(Site), levels = levels(Bream$Site))  
  }) %>% arrange(Site)  
Scores
```

Site	Score	11-12	12-13	13-14	14-15	15-16	16-17
Ramsay Crossing	0.84					0.58	0.62
Munduran Creek	0.31	0.81	0.24	0.33	0.49	0.37	0.71
Black Swan	0.68				0.67	0.06	0.97
Targinnie Creek	0.53	0.67	0.18		0.87	0.16	0.74
Graham Creek	0.31				0.77	0.31	0.56
Hobble Gully	0.37				0.56	0.49	0.62
Mud Island	0.16					0.17	0.79
Boat Creek	0.09		0.21	0.43	0.80	0.28	0.47
Little Enfield Creek	0.60				0.76	0.28	0.71
Barney Point Pond	0.05		0.22	0.44	0.70	0.24	0.57
Beecher Creek	0.51	0.84	0.12	0.40	0.66	0.29	0.76
Old Bruce Highway Bridge	0.35				0.45	0.35	0.83
Callemondah	0.78	0.61	0.08	0.39	0.68	0.49	0.91
Farmers Point	0.25					0.09	0.91
Gatcombe Anchorage	0.06					0.34	0.53
Wappentake Creek	0.26		0.22	0.49	0.59	0.32	0.67
South Trees	0.48					0.42	0.64
Crematorium Pool	0.77					0.27	0.85
Old Boyne	0.86	0.69	0.30		0.61	0.44	0.71
Boyne Highway	0.81				0.54	0.49	0.78
Broadacres	0.60					0.36	0.73
Iveragh	0.62					0.41	0.70
Oaky	0.93					0.49	0.76
7 Mile	0.94					0.50	0.75
Worthington	0.19					0.27	0.69
Sandy Bridge	0.87					0.39	0.80

Notes

- The offset of $\log(\text{Casts})$ allows for minor fluctuations from the nominal 20 casts per visit and recognises that, other things being equal, catch should be proportional to effort.
- The `Month` fixed effect allows for systematic changes in productivity over the calendar year.
- The `Site` score is an estimate of the (relative) productivity of the site for Bream recruitment overall, allowing for the minimal effect of `Depth` and `Rock` in the fixed effects.
- The `Year` \times `Site` scores give a measure of the *change* in catch rate from year to year *relative* to what might have been expected given the productivity of the site.
- Converting scores to a (0, 1) scale is something required by the Health Card protocol; this is only *one possible* way to do it.

A Theme functions

ggplot2 release 2.0.0 (or later) has a wide range of pre-defined themes, as well as theme generating and theme modifying functions. A list is:

```
apropos("^ (theme$|theme_)") %>% noquote
```

[1] theme	theme_base	theme_bw
[4] theme_calc	theme_classic	theme_clean
[7] theme_dark	theme_economist	theme_economist_white
[10] theme_excel	theme_excel_new	theme_few
[13] theme_fivethirtyeight	theme_foundation	theme_gdocs
[16] theme_get	theme_gray	theme_grey
[19] theme_hc	theme_igray	theme_light
[22] theme_linedraw	theme_map	theme_minimal
[25] theme_pander	theme_par	theme_replace
[28] theme_set	theme_solarized	theme_solarized_2
[31] theme_solid	theme_stata	theme_test
[34] theme_tufte	theme_update	theme_void
[37] theme_wsj		

B Two helper functions

These are needed to define harmonic terms and interactions.

```
Harm <- function (theta, k = 4) {  
  X <- matrix(0, length(theta), 2 * k)  
  nam <- as.vector(outer(c("c", "s"), 1:k, paste, sep = ""))  
  dimnames(X) <- list(names(theta), nam)  
  m <- 0  
  for (j in 1:k) {  
    X[, (m <- m + 1)] <- cos(j * theta)  
    X[, (m <- m + 1)] <- sin(j * theta)  
  }  
  X  
}  
  
Hyear <- function(x, k = 4)  
  Harm(2*base::pi*x/365.25, k)  
  
twoWay <- local({  
  `%star%` <- function(X, Y) { ## all column-products  
    X <- as.matrix(X)  
    Y <- as.matrix(Y)  
    stopifnot(is.numeric(X), is.numeric(Y),
```

```

        nrow(X) == nrow(Y))
XY <- matrix(NA, nrow(X), ncol(X)*ncol(Y))
k <- 0
for(i in 1:ncol(X))
  for(j in 1:ncol(Y)) {
    k <- k+1
    XY[, k] <- X[, i] * Y[, j]
  }
  XY
}
function(day, sea, k = c(3,2)) {
  Hyear(day, k[1]) %star% splines::ns(sea, k[2])
}
})
Store(Harm, Hyear, twoWay, lib = .Robjcts)

```

References

- Cribari-Neto, F. and A. Zeileis (2010). Beta regression in **R**. *Journal of Statistical Software* 34(2), 1–24.
- Venables, W. N. and B. D. Ripley (2002). *Modern Applied Statistics with S* (Fourth ed.). New York: Springer. ISBN 0-387-95457-0.
- Wood, S. (2011). *gamm4: Generalized additive mixed models using mgcv and lme4*. CRAN. R package version 0.1–5.

Session information

Date: 2021-01-29

- R version 4.0.3 (2020-10-10), x86_64-pc-linux-gnu
- Running under: Ubuntu 20.04.1 LTS
- Matrix products: default
- BLAS: /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.9.0
- LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.9.0
- Base packages: base, datasets, graphics, grDevices, methods, stats, utils
- Other packages: dplyr 1.0.3, english 1.2-5, forcats 0.5.1, ggplot2 3.3.3, ggthemes 4.2.4, gridExtra 2.3, knitr 1.31, lattice 0.20-41, lme4 1.1-26, Matrix 1.3-2, patchwork 1.1.1, purrr 0.3.4, readr 1.4.0, scales 1.1.1, SOAR 0.99-11, stringr 1.4.0, tibble 3.0.5, tidyr 1.1.2, tidyverse 1.3.0, WWRCourse 0.2.3, WWRData 0.1.0, WWRGraphics 0.1.2, WWRUtilities 0.1.2, xtable 1.8-4
- Loaded via a namespace (and not attached): assertthat 0.2.1, backports 1.2.1, boot 1.3-26, broom 0.7.3, cellranger 1.1.0, cli 2.2.0, colorspace 2.0-0, compiler 4.0.3, crayon 1.3.4, DBI 1.1.1, dbplyr 2.0.0, digest 0.6.27, ellipsis 0.3.1, evaluate 0.14, fansi 0.4.2, farver 2.0.3, fractional 0.1.3, fs 1.5.0, generics 0.1.0, ggrepel 0.9.1, glue 1.4.2, grid 4.0.3, gtable 0.3.0, haven 2.3.1, highr 0.8, hms 1.0.0, http 1.4.2, iterators 1.0.13, jsonlite 1.7.2, labeling 0.4.2, lazyData 1.1.0, lifecycle 0.2.0, lubridate 1.7.9.2, magrittr 2.0.1, MASS 7.3-53, mgcv 1.8-33, minqa 1.2.4, modelr 0.1.8, munsell 0.5.0, nlme 3.1-151, nloptr 1.2.2.2, parallel 4.0.3, PBSmapping 2.73.0, pillar 1.4.7, pkgconfig 2.0.3, R6 2.5.0, randomForest 4.6-14, Rcpp 1.0.6, readxl 1.3.1, reprex 1.0.0, rlang 0.4.10, rpart 4.1-15, rstudioapi 0.13, rvest 0.3.6, splines 4.0.3, statmod 1.4.35, stringi 1.5.3, tidyselect 1.1.0, tools 4.0.3, vctrs 0.3.6, withr 2.4.1, xfun 0.20, xml2 1.3.2