

Working with 

A University of Queensland Advanced Workshop

Session 9: Machine Learning Approaches

Bill Venables, CSIRO/Data61, Dutton Park

Rhetta Chappell, Griffith University

2–5 February, 2021

Contents

1	The Swiss credit card data set	3
2	Trees and forests	7
3	Do you want a credit card?	9
3.1	An initial tree model	9
3.2	Simple bagging	16
3.3	A parallel version	21
3.4	The actual random forest	24
3.5	Progressive error rate	26
3.6	Variable importances	30
3.6.1	Partial plots for predictor variables	33
3.7	Parametric models	35

3.8	Boosting other models	37
3.9	Weird science: C5.0	38
3.10	Gradient boosting models, GBM	41
4	The final reckoning	46
	References	49
	Session information	50

1 The Swiss credit card data set

For this section, we use one data set which we prepare from original sources, namely the file `creditCard.csv`.

- Response: *credit.card.owner*, binary with levels `c("no", "yes")`.
- Predictors: mostly numeric, to do with banking activity.
 - *profession* is a sparse factor which we re-group into three densely populated levels (as factor *pclass*).
 - *nationality* which we re-group into two levels, *"Swiss"* and *"Foreigner"*.
 - *sex* is the only other factor.
- Entries have a numerical *ident* which we zero-fill and move to the *rownames* rather than have it as a potentially active variable.

Read in the data and

```
library(WWRCourse)
CCf <- system.file("extdata", "creditCard.csv", package = "WWRCourse")
CC <- read.csv(CCf, na.strings="", stringsAsFactors = TRUE)
row.names(CC) <- fill0(CC$ident)

with(CC, c(cases = nrow(CC), known = sum(!is.na(credit.card.owner))))

cases known
10893  2085
```

The data is mostly *unknown*: for building predictive models we only have two thousand and eighty-five *at most* from ten thousand eight hundred and ninety-three cases.

```

CC <- CC %>% within({
  p2 <- c("doctor", "engineer", "lawyer", "professor", "business")
  p1 <- c("teacher", "police", "service", "chemist", "nurse",
          "postman", "physical")
  p0 <- "none"
  pclass <- rep("", length(profession))
  pclass[profession %in% p0] <- "p0"
  pclass[profession %in% p1] <- "p1"
  pclass[profession %in% p2] <- "p2"
  pclass <- factor(pclass, levels = paste0("p", 0:2))
  swiss <- ifelse(nationality == "CH", "Swiss", "Foreigner")
  sex <- factor(sex)
  swiss <- factor(swiss)
  pclass <- factor(pclass)
  ident <- profession <- nationality <- p0 <- p1 <- p2 <- NULL
})
CCKnown <- na.omit(CC)
dim(CCKnown)

[1] 1620  55

```

For model building we have just one thousand six hundred and twenty

cases and fifty-four potential predictors.

For demonstration purposes we split these into *Training* and Test data sets.

```
set.seed(20200202)    ## for reproducibility
train <- sample(nrow(CCKnown), nrow(CCKnown)/2)

CCTrain <- CCKnown[train, ]
CCTest <- CCKnown[-train, ]
Store(CC, CCKnown, CCTrain, CCTest, train, remove = FALSE) ## for safety
```

2 Trees and forests

- A technique that developed in machine learning and now widely used in data mining.
- The model uses *recursive partitioning* of the data and is a greedy algorithm.
- The two main types of tree models are
 - ***Regression trees*** — response is a continuous variable and fitting uses a least squares criterion,
 - ***Classification trees*** — response is a factor variable and fitting uses an entropy (multinomial likelihood) criterion.

- Model fitting is easy. Inference poses more of a dilemma.
- The tree structure is very unstable. *boosting* and *bagging* (random forests) can be useful ways around this.
- Two packages for tree models: `rpart` (which is part of **R** itself) and the older `tree`, (Ripley., 2012), which has an **S-PLUS** flavour and a few advantages for teaching.

Use `rpart` in practice.

3 Do you want a credit card?

Our plan is to build predictive models on the *Training* set, and use the *Test* set to see how well they fared. We employ a number of techniques, mostly of the “black box” kind.

3.1 An initial tree model

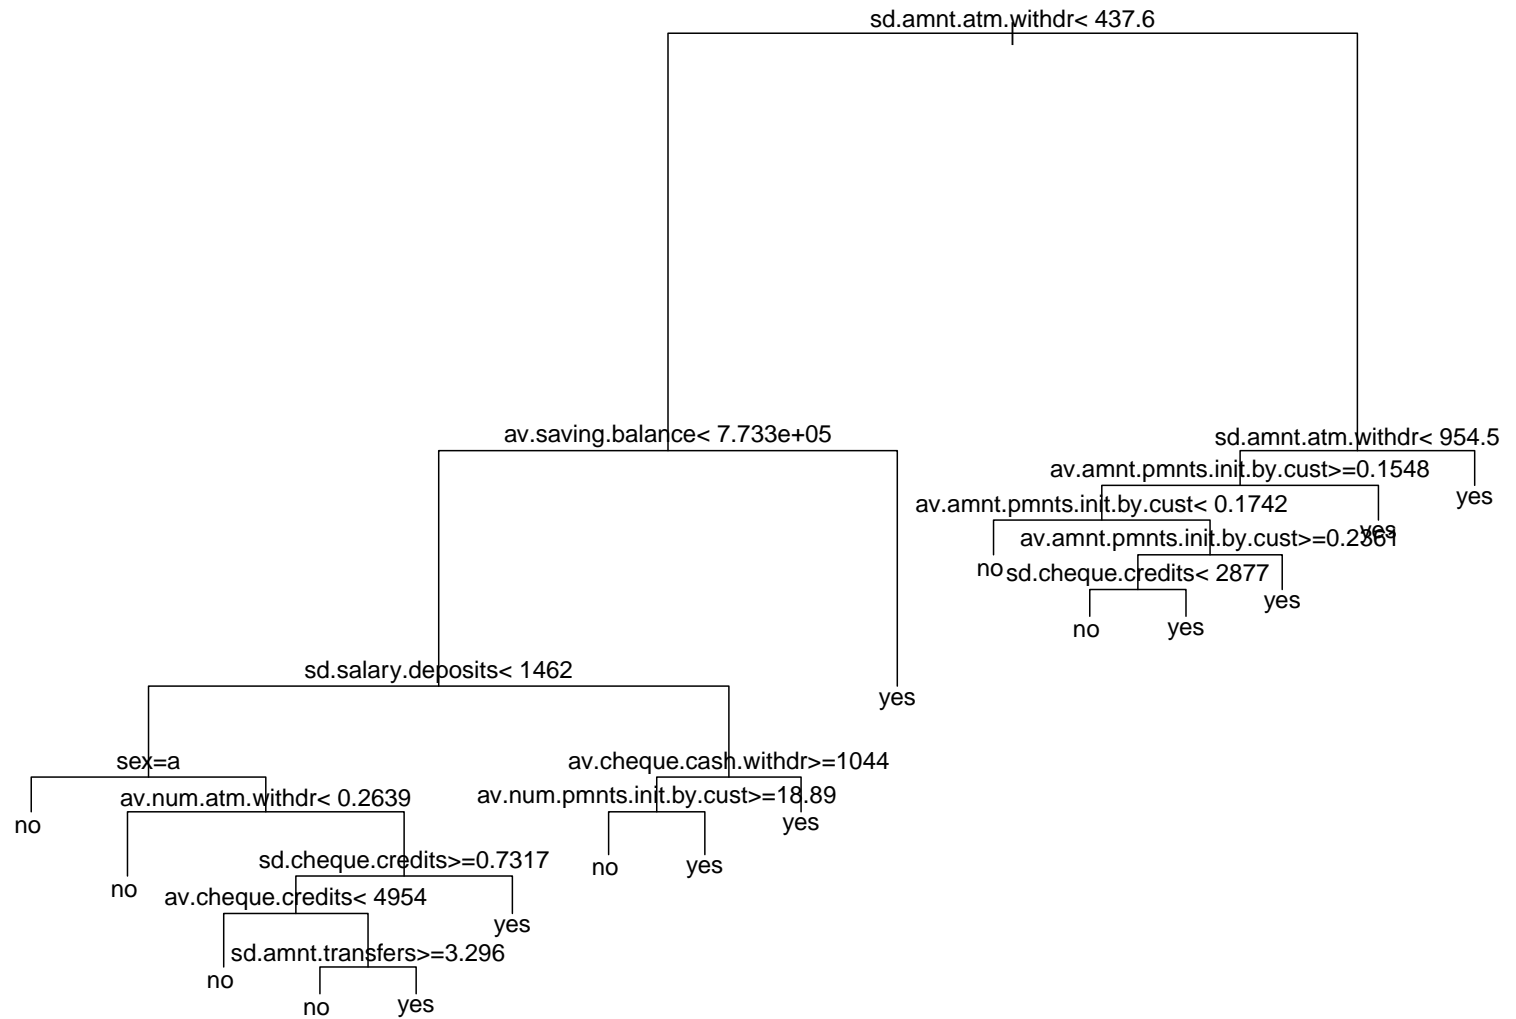
```
requireData(rpart)
CCTree <- rpart(credit.card.owner ~ ., CCTrain)
Store(CCTree, remove = FALSE)

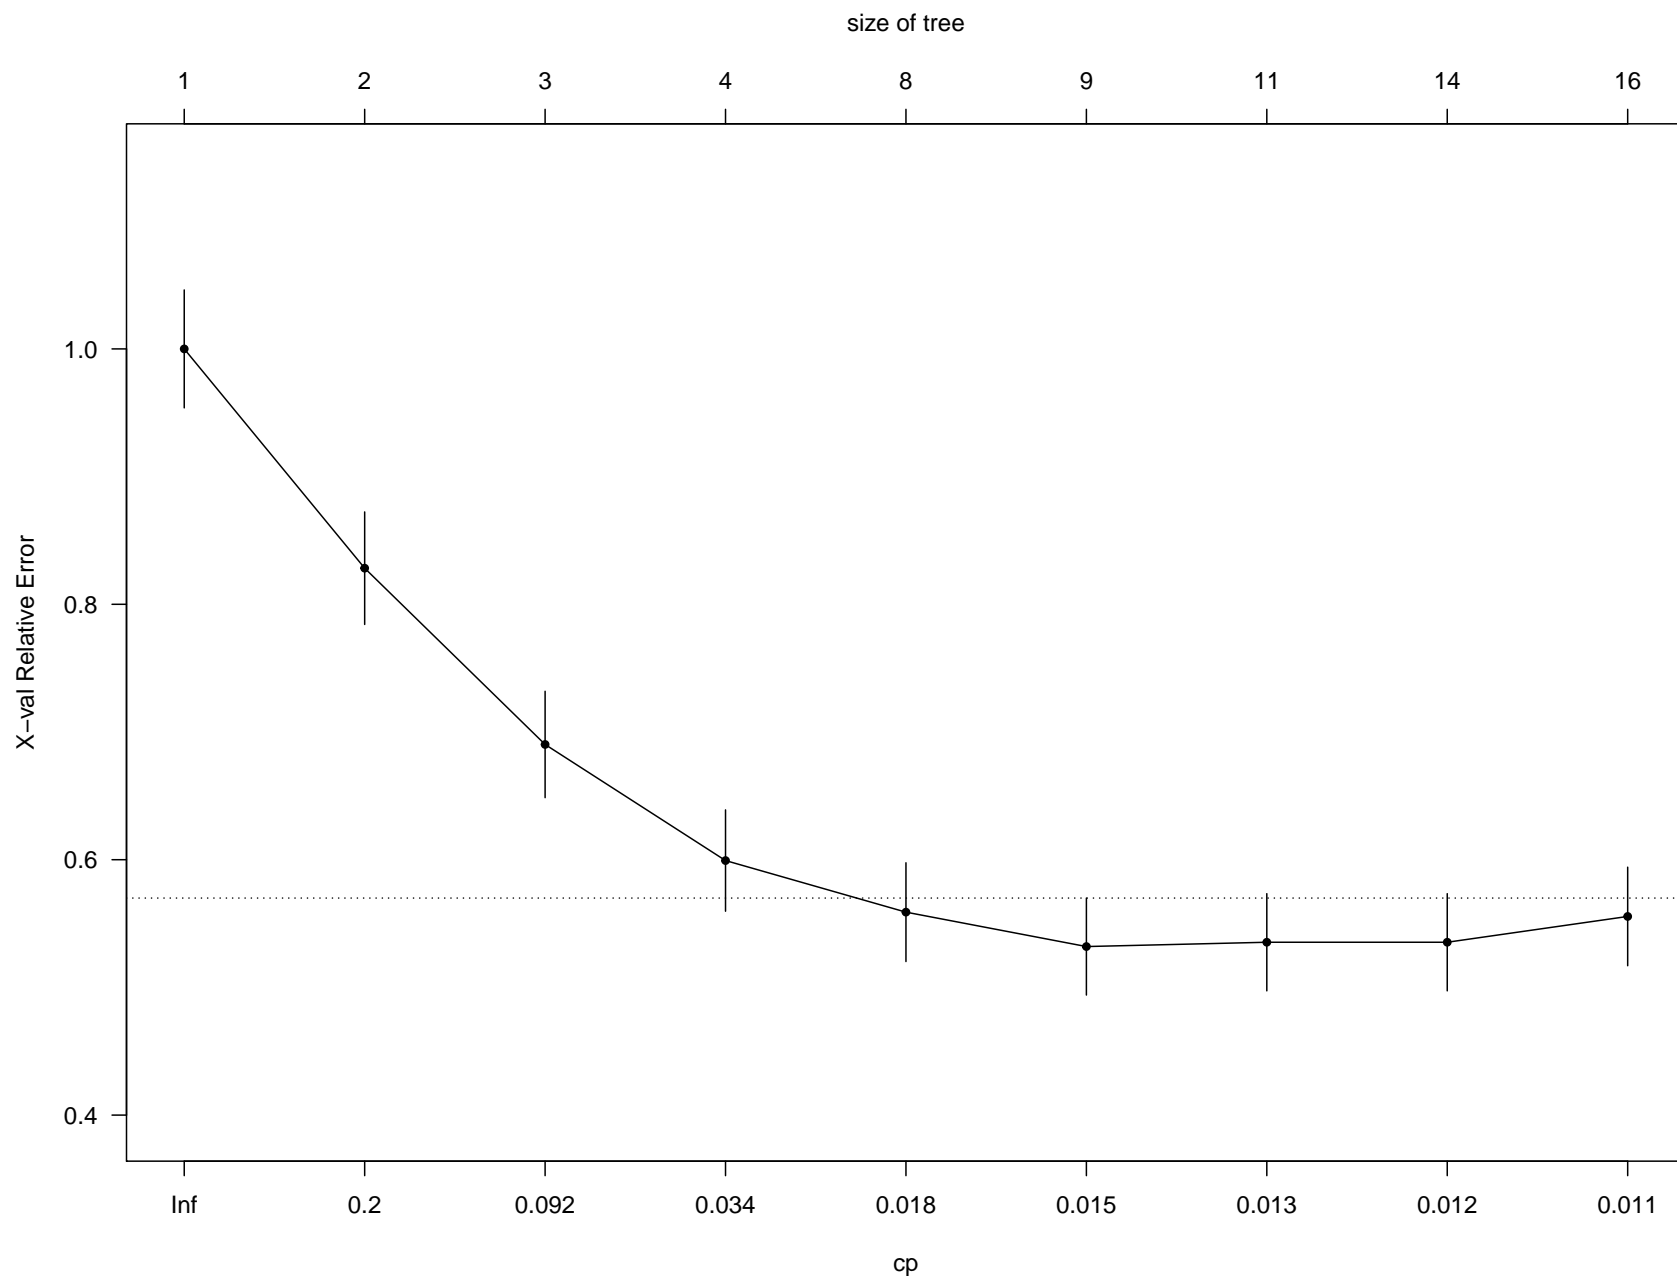
## The first graphic shows the initial fitted tree:

plot(CCTree)
text(CCTree, xpd = NA)

## The next graphic is to check for the need to prune:

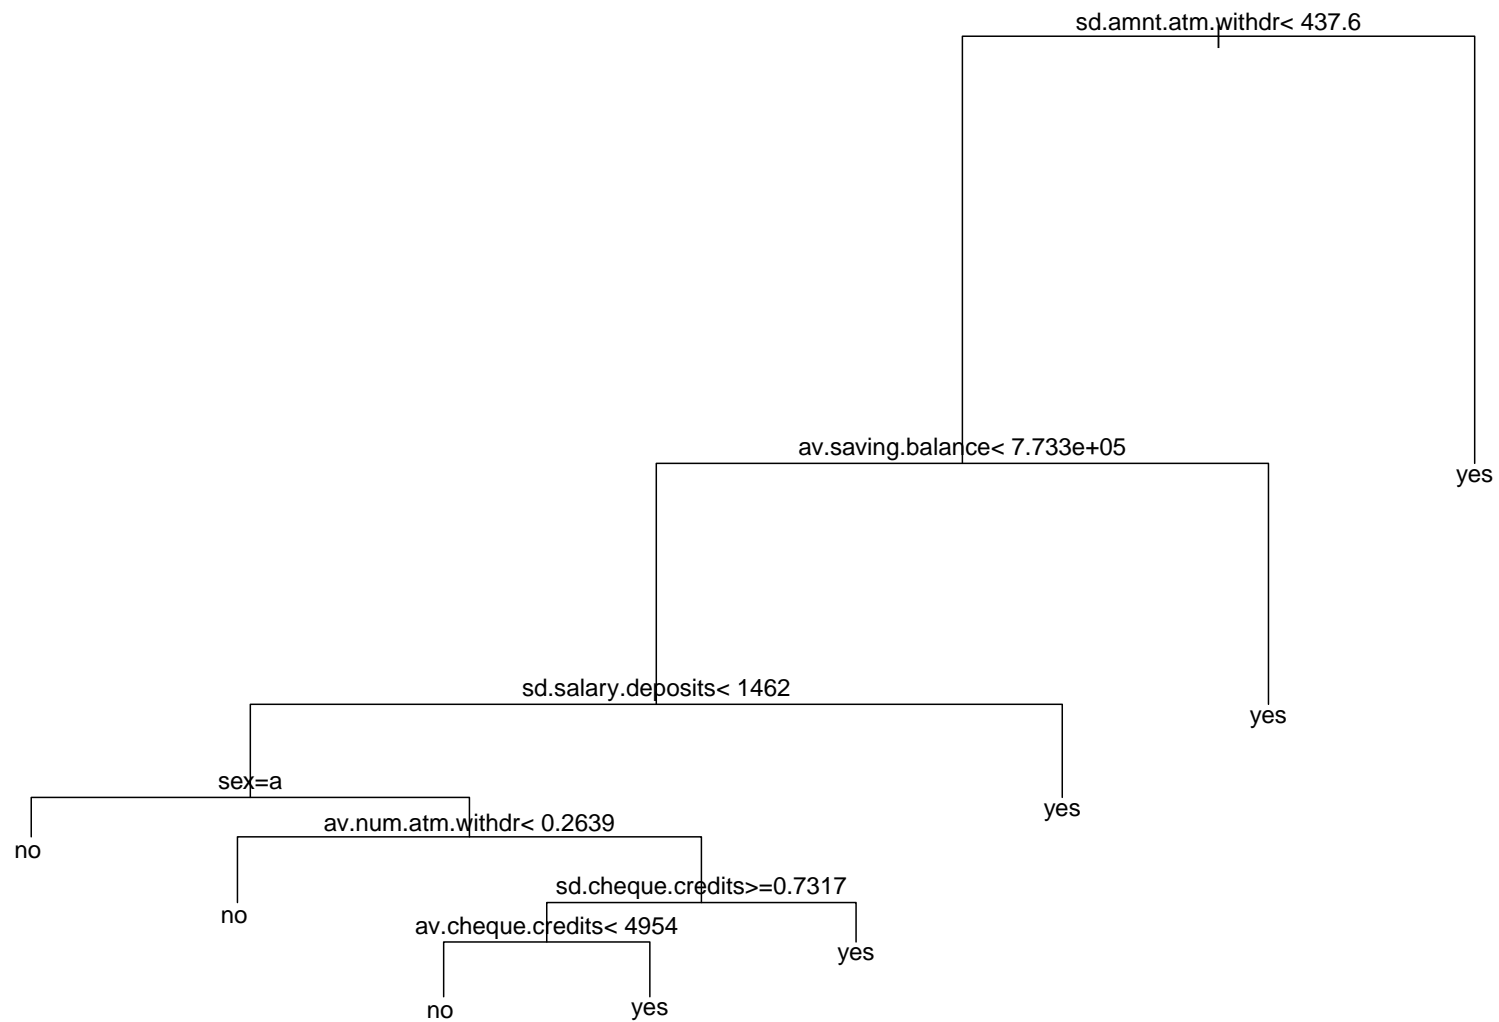
plotcp(CCTree)
```





Pruning is suggested by the “one standard error” rule. Get the pruned tree:

```
oneSERule(CCTree)                ## optimal tuning parameter  
[1] 0.01683502  
  
CCPTree <- prune(CCTree, cp = oneSERule(CCTree))  
Store(CCPTree, remove = FALSE)  
  
plot(CCPTree)                    ## should have 10 terminal nodes  
text(CCPTree, xpd = NA)
```



The “one standard error rule” function(s) are listed here for completeness. (The coding details are not important for our primary purpose here.)

```
WWRCourse::oneSERule  
  
function (tree, f, ...)  
{  
  UseMethod("oneSERule")  
}  
<bytecode: 0x5641497423b8>  
<environment: namespace:WWRCourse>  
  
methods("oneSERule")  
  
[1] oneSERule.rpart*  
see '?methods' for accessing help and source code  
  
WWRCourse:::oneSERule.rpart  
  
function (tree, f = 1, ...)  
{  
  cp <- data.frame(tree$cptable)  
  imin <- with(cp, which(xerror == min(xerror))[1])  
  with(cp, CP[which(xerror <= xerror[imin] + f * xstd[imin])[1]])  
}  
<bytecode: 0x56414d4708c0>
```

```
<environment: namespace:WWRCourse>
```


3.2 Simple bagging

“Bootstrap aggregation” — invented by Leo Breimann as a device to stabilize tree methods and improve their predictive capacity. Very much a “black box” technique.

- Grow a forest of trees using bootstrap samples of the training data.
- For predictions average over the forest:
 - For classification trees, take a majority vote,
 - For regression trees, take an average.

‘Random forests’, (Liaw and Wiener, 2002), is a development of bagging with extra protocols imposed.

Consider bagging “by hand”.

```
bagRpart <- local({
  bsample <- function(dataFrame) # bootstrap sampling
    dataFrame[sample(nrow(dataFrame), rep = TRUE), ]
  function(object, data = eval.parent(object$call$data),
    nBags=200, type = c("standard", "bayesian"), ...) {
    type <- match.arg(type)
    bagsFull <- vector("list", nBags)
    if(type == "standard") {
      for(j in 1:nBags)
        bagsFull[[j]] <- update(object, data = bsample(data))
    } else {
      nCases <- nrow(data)
      for(j in 1:nBags)
        bagsFull[[j]] <- update(object, data = data, weights = rexp(nCases))
    }
    class(bagsFull) <- "bagRpart"
    bagsFull
  }
})

## a prediction method for the objects (somewhat tricky!)
```

```
predict.bagRpart <- function(object, newdata, ...) {  
  X <- sapply(object, predict, newdata = newdata, type = "class")  
  candidates <- levels(predict(object[[1]], type = "class"))  
  X <- t(apply(X, 1, function(r) table(factor(r, levels = candidates))))  
  factor(candidates[max.col(X)], levels = candidates)  
}  
Store(bagRpart, lib = .Robjects)
```

An alternative coding using *replicate*:

```

bagRpart <- local({
###
  bsample <- function(dataFrame) # bootstrap sampling
    dataFrame[sample(nrow(dataFrame), rep = TRUE), ]
###
  function(object, data = eval.parent(object$call$data),
            nBags=200, type = c("standard", "bayesian"), ...) {
    type <- match.arg(type)
    bagsFull <- if(type == "standard") {
      replicate(nBags, update(object, data = bsample(data)),
                simplify = FALSE)
    } else {
      nCases <- nrow(data)
      replicate(nBags, update(object, data = data, weights = rexp(nCases)),
                simplify = FALSE)
    }
    class(bagsFull) <- "bagRpart"
    bagsFull
  }
})

```

Now for an object or two:

```
set.seed(4321) ##
Obj <- update(CCTree, cp = 0.005, minsplit = 9) ## expand the tree
(one <- rbind(standard = system.time(CCSBag <- bagRpart(Obj, nBags = 200)),
             Bayes = system.time(CCBBag <- bagRpart(Obj, nBags = 200,
                                                    type = "bayes")))))

      user.self sys.self elapsed user.child sys.child
standard    11.333    0.008   11.342         0         0
Bayes       11.561    0.004   11.565         0         0

Store(CCSBag, CCBBag, remove = FALSE)
```

3.3 A parallel version

Some preliminaries first:

```
## parallel backend; includes other pkgs
suppressPackageStartupMessages({
  library(doParallel)
})

(nc <- detectCores()) ## how many CPUs has your computer?

[1] 12

cl <- makeCluster(nc-1)
registerDoParallel(cl)
```

A modified version of the fitting function. Some care needed:

```
bagRpartParallel <- local({
  bsample <- function(dataFrame) # bootstrap sampling
    dataFrame[sample(nrow(dataFrame), rep = TRUE), ]

  function(object, data = eval.parent(object$call$data),
    nBags = 200, type = c("standard", "bayesian"), ...,
    cores = detectCores() - 1, seed0 = as.numeric(Sys.Date())) {
    type <- match.arg(type)
    bagsFull <- foreach(j = idiv(nBags, chunks=cores), seed = seed0+seq(cores),
      .combine = c, .packages = c("rpart", "stats"),
      .inorder = FALSE, .export = c("bsample")) %dopar% {
      ## now inside a single core
      set.seed(seed = seed)
      if(type == "standard") {
        replicate(j, simplify = FALSE,
          update(object, data = bsample(data)))
      } else {
        replicate(j, simplify = FALSE,
          update(object, data = data,
            weights = rexp(nrow(data))))
      }
    }
  }
})
```

```

                                ## end of single core mode
                                }
class(bagsFull) <- "bagRpart"
bagsFull
}
})

```

A timing comparison

```

Obj <- update(CCTree, cp = 0.005, minsplit = 9)  ## expand the tree
rbind(one,
      standardP = system.time(CCSBagP <- bagRpartParallel(Obj, nBags = 200)),
      BayesP = system.time(CCBBagP <- bagRpartParallel(Obj, nBags = 200,
                                                         type = "bayes")))

```

	user.self	sys.self	elapsed	user.child	sys.child
standard	11.333	0.008	11.342	0.000	0
Bayes	11.561	0.004	11.565	0.000	0
standardP	0.156	0.016	1.993	0.002	0
BayesP	0.152	0.008	1.924	0.003	0

```

rm(Obj, one)
Store(CCSBagP, CCBBagP, remove = FALSE)

```


3.4 The actual random forest

The random forest package, (Liaw and Wiener, 2002), implements this technology, and more, automatically. The number of trees is set to 500 by default. How many times does each observation get sampled if we restrict it to 200 trees?

```
n <- nrow(CCTrain)
X <- replicate(200,
               table(factor(sample(n, rep=TRUE), levels = 1:n)))
(lims <- range(rowSums(X > 0)))

[1] 104 150

rm(n, X)
```

So in this simulation the cases were sampled between one hundred and four and one hundred and fifty times. This seems about enough.

We now fit the random forest.

```
requireData(randomForest)
(CCRf <- randomForest(credit.card.owner ~ ., CCTrain, ntree = 200))

Call:
randomForest(formula = credit.card.owner ~ ., data = CCTrain,          ntree = 200)
      Type of random forest: classification
      Number of trees: 200
No. of variables tried at each split: 7

      OOB estimate of  error rate: 14.69%
Confusion matrix:
      no yes class.error
no  217  80  0.26936027
yes   39 474  0.07602339

Store(CCRf, remove = FALSE)
```


The “out of bag” (OOB) error rate is estimated as 14.69%. We check later how this compares with the observed error rate in the reserved *Test* data.

3.5 Progressive error rate

The random forest technology keeps a tally of the OOB error rate as the forest grows.

```
err <- as.data.frame(CCRf$err.rate) %>%  
  rownames_to_column("trees") %>%  
  mutate(trees = as.numeric(trees)) %>%  
  pivot_longer(cols = -trees, names_to = "Which", values_to = "Error")  
ggplot(err) + aes(x = trees, y = Error, colour = Which) + geom_line() +  
  scale_colour_brewer(palette = "Dark2", name = "Which one:") +  
  ylim(0, max(err$Error)*1.05) + labs(title = "Credit Card - Training Data") +  
  guides(colour = guide_legend(ncol = 3)) +  
  theme_bw() + theme(legend.position = c(0, 0)+0.01,  
                     legend.justification = c("left", "bottom"),  
                     plot.title = element_text(hjust = 0.5, face = "bold"))
```


Credit Card – Training Data



We can do this more directly with the *plot* method for random forest objects. As is our custom, we add a few little aesthetic enhancements:

```
ER <- CCRf$err.rate
plot(CCRf, ylim = range(0, ER)*1.05, lty = "solid", las = 1,
      panel.first = grid(lty = "dashed"), main = "Credit Card - Training Data")
legend("bottomleft", colnames(ER), ncol = ncol(ER), bg = "white",
      col = 1:3, lty = "solid", bty = "o", box.col = "transparent",
      title = "Which one:")
```

Credit Card – Training Data




3.6 Variable importances

One other nice by-product is variable importances.

```
par(family="sans")  
varImpPlot(CCRf, pch = 20, col = "navy") ## causes a plot
```

CCRF

av.amnt.atm.withdr
 av.salary.deposits
 av.cheque.credits
 sd.amnt.atm.withdr
 sd.cheque.credits
 av.saving.balance
 av.cheque.debits
 age
 sd.amnt.transfers
 av.cheque.cash.withdr
 sd.cheque.debits
 av.num.atm.withdr
 sd.salary.deposits
 sd.cheque.cash.withdr
 sex
 sd.num.atm.withdr
 sd.amnt.pmnts.init.by.cust
 av.amnt.transfers
 av.num.cheque.cash.withdr
 av.num.pmnts.init.by.cust
 sd.num.pmnts.init.by.cust
 av.amnt.pmnts.init.by.cust
 sd.saving.balance
 av.num.salary.deposits
 sd.num.cheque.cash.withdr
 sd.num.saving.cash.deposits
 av.num.saving.cash.deposits
 av.amnt.saving.cash.deposits
 sd.amnt.saving.cash.deposits
 pclass




```
(v <- sort(drop(importance(CCRf)), decreasing = TRUE))[1:6]
```

av.amnt.atm.withdr	av.salary.deposits	av.cheque.credits
26.12861	25.72058	25.68204
sd.amnt.atm.withdr	sd.cheque.credits	av.saving.balance
25.26308	19.25392	17.21673

```
best_few <- names(v)[1:20] %>% print ## used later
```


[1] "av.amnt.atm.withdr"	"av.salary.deposits"
[3] "av.cheque.credits"	"sd.amnt.atm.withdr"
[5] "sd.cheque.credits"	"av.saving.balance"
[7] "av.cheque.debits"	"age"
[9] "sd.amnt.transfers"	"av.cheque.cash.withdr"
[11] "sd.cheque.debits"	"av.num.atm.withdr"
[13] "sd.salary.deposits"	"sd.cheque.cash.withdr"
[15] "sex"	"sd.num.atm.withdr"
[17] "sd.amnt.pmnts.init.by.cust"	"av.amnt.transfers"
[19] "av.num.cheque.cash.withdr"	"av.num.pmnts.init.by.cust"

3.6.1 Partial plots for predictor variables

These look at the way predictors appear to influence the response, *mutatis mutandis*.

```
partialPlot(CCRf, pred.data = CCTrain, x.var = best_few[1], xlab = best_few[1])
```

Partial Dependence on best_few[1]



3.7 Parametric models

Tree models and random forests are natural competitors to the standard parametric models, notably GLMs. We begin with a naive model based only on what appear good variables in the random forest, and then consider other modest versions, but automatically produced.

```
(form <- as.formula(paste("credit.card.owner~",
                          paste(best_few, collapse="+"))))

credit.card.owner ~ av.amnt.atm.withdr + av.salary.deposits +
  av.cheque.credits + sd.amnt.atm.withdr + sd.cheque.credits +
  av.saving.balance + av.cheque.debits + age + sd.amnt.transfers +
  av.cheque.cash.withdr + sd.cheque.debits + av.num.atm.withdr +
  sd.salary.deposits + sd.cheque.cash.withdr + sex + sd.num.atm.withdr +
  sd.amnt.pmnts.init.by.cust + av.amnt.transfers + av.num.cheque.cash.withdr +
  av.num.pmnts.init.by.cust

Call <- substitute(glm(FORM, binomial, CCTrain), list(FORM = form))
CCglmNaive <- eval(Call)
Store(CCglmNaive, remove = FALSE)
```

```

upp <- paste("~", paste(setdiff(names(CCTrain),
                                "credit.card.owner"),
                        collapse="+")) %>%
  as.formula()
scope <- list(upper=upp, lower=~1)
CCGlmAIC <- step_AIC(CCGlmNaive, scope = scope)
CCGlmGIC <- step_GIC(CCGlmNaive, scope = scope)
CCGlmBIC <- step_BIC(CCGlmNaive, scope = scope)
Store(CCGlmAIC, CCGlmBIC, form, upp, remove = FALSE)

```

3.8 Boosting other models

The *mboost* package is an implementation of the boosting idea for a suite of different models. It is described in [Buehlmann and Hothorn \(2007\)](#). We consider two different kinds of boosted models, namely a boosted GLM and a boosted TREE version.

The package is profligate in the number of others it requires!

```
requireData(mboost)
Call <- substitute(glmboost(FORM, data = CCTrain, family=Binomial()),
                    list(FORM = form)) ## same as naive model
CCglmboost <- eval(Call)

## more packages needed!
CCblackboost <- blackboost(credit.card.owner ~ ., CCTrain,
                           family = Binomial())
Store(CCglmboost, CCblackboost, remove = FALSE)
```

3.9 Weird science: C5.0

This is a development of a previously proprietary algorithm of [Quinlan \(1993\)](#); [Kuhn and Quinlan \(2020\)](#). It yields a rule-based classifier that might be considered intermediate between trees and forests.

```
requireData(C50)  ## NB C-5-Zero
(CCc50 <- C5.0(credit.card.owner ~ ., CCTrain))
```

Call:

```
C5.0.formula(formula = credit.card.owner ~ ., data = CCTrain)
```

Classification Tree

Number of samples: 810

Number of predictors: 54

Tree size: 37

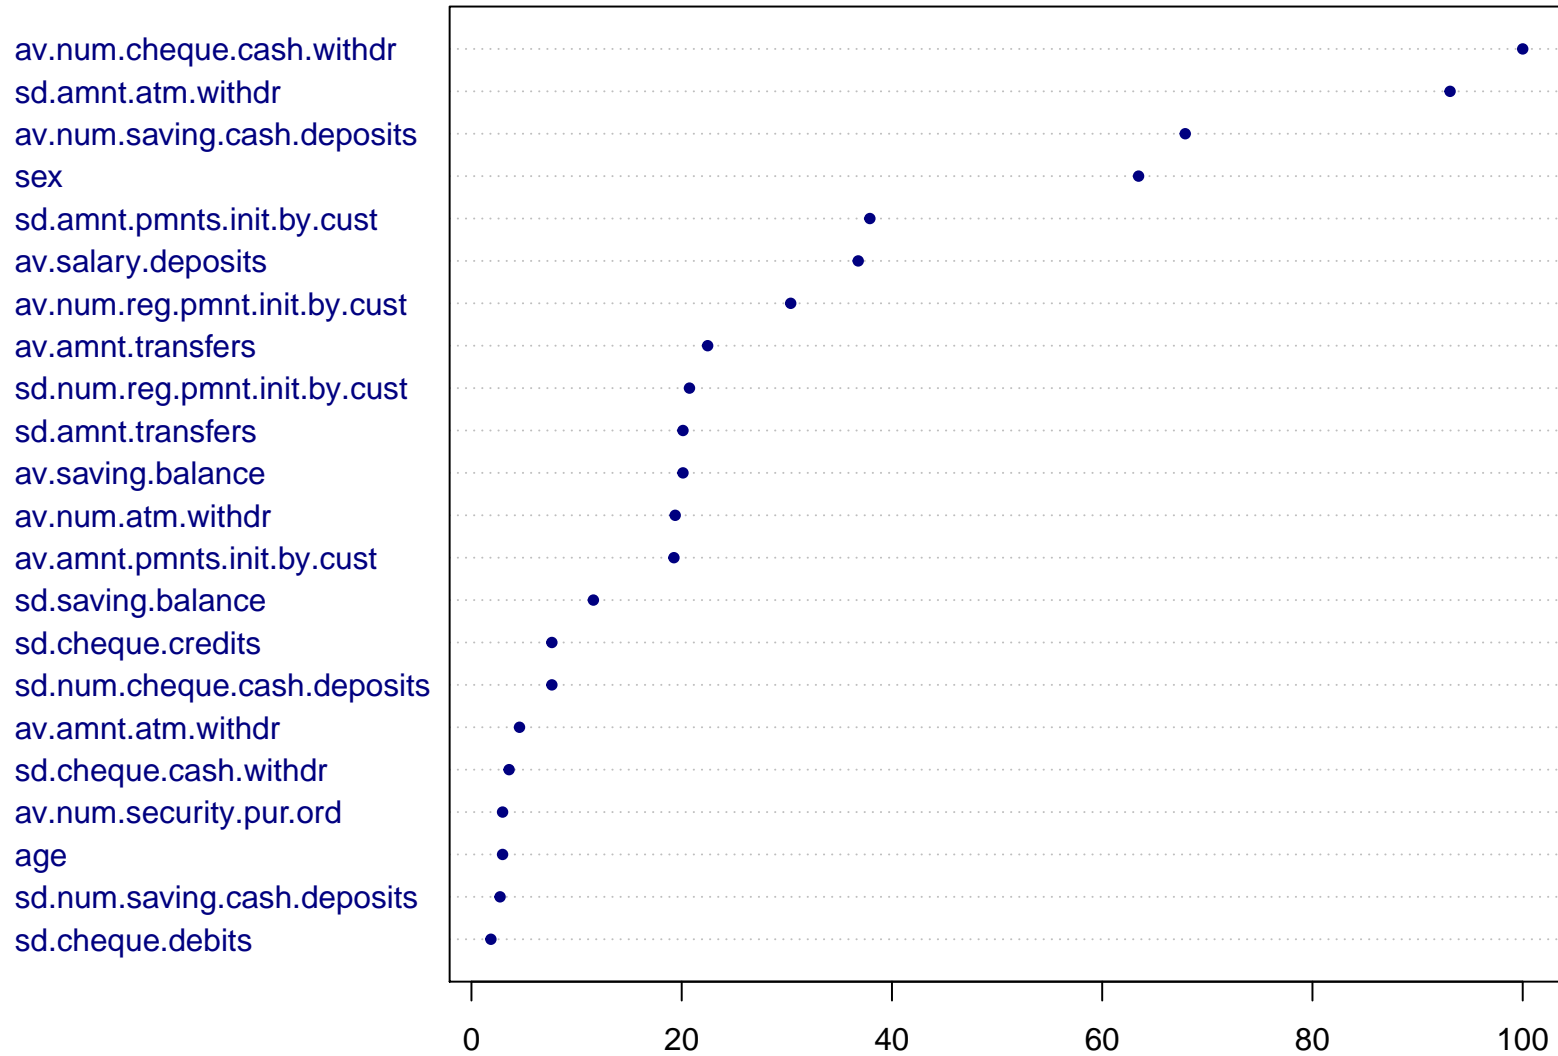
Non-standard options: attempt to group attributes

```
Store(CCc50)
```

The `C50` package, like `randomForest`, has some tools that allow the fitted model object to be investigated.

These include tools to assess *variable importance*, where the measure itself is somewhat obscure, but it expressed as a percentage.

```
obj <- C5imp(CCc50)
data.frame(importance = obj$Overall, variable = rownames(obj)) %>%
  filter(importance > 0) %>%
  arrange(importance) %>%
  with(., dotchart(importance, as.character(variable), pch = 20, col = "navy"))
rm(obj)
```

3.10 Gradient boosting models, GBM

Like [C5.0](#) this method was proprietary but for some time now there has been an [R](#) implementation by Greg Ridgeway, ([Greenwell et al., 2020](#)).

It offers a powerful machine learning method for several classes of models, including 'bernoulli' (binary data, but **not** multinomial, unlike [randomForest](#) or even [C5.0](#).)

Bernoulli models need to be fitted with a numeric response confined to the values $\{0, 1\}$, which can be a bit awkward:

```
suppressPackageStartupMessages(library(gbm))
(CCgbm <- gbm(ifelse(credit.card.owner == "no", 0, 1) ~ ., data = CCKnown,
              distribution = "bernoulli", n.trees = 200, cv.folds = 10))

gbm(formula = ifelse(credit.card.owner == "no", 0, 1) ~ ., distribution = "bernoulli",
     data = CCKnown, n.trees = 200, cv.folds = 10)
```

A gradient boosted model with bernoulli loss function.

200 iterations were performed.

The best cross-validation iteration was 200.

There were 54 predictors of which 34 had non-zero influence.


There are several useful tools for investigating the fit and efficacy of a GBM. The first of these is one for looking at the relative influence of the predictors:

```
summary(CCgbm, plotit = FALSE) %>% filter(rel.inf > 0)
```

	var	rel.inf
sd.amnt.atm.withdr	sd.amnt.atm.withdr	20.2703539
sex	sex	11.2280381
av.amnt.atm.withdr	av.amnt.atm.withdr	11.0621992
av.cheque.credits	av.cheque.credits	9.5013640
av.salary.deposits	av.salary.deposits	5.6615081
av.saving.balance	av.saving.balance	4.6562862
av.num.cheque.cash.withdr	av.num.cheque.cash.withdr	4.5393055
....		
av.num.saving.cash.withdr	av.num.saving.cash.withdr	0.2743822
av.cheque.cash.withdr	av.cheque.cash.withdr	0.2336712
sd.num.cheque.cash.withdr	sd.num.cheque.cash.withdr	0.2199467
av.amnt.reg.pmnt.init.by.cust	av.amnt.reg.pmnt.init.by.cust	0.1665780
sd.num.security.sales.ord	sd.num.security.sales.ord	0.1537358
av.amnt.pmnts.init.by.cust	av.amnt.pmnts.init.by.cust	0.1441431
av.num.security.sales.ord	av.num.security.sales.ord	0.1435577

But wait, there's more. You can get a *barplot* of the result using same function

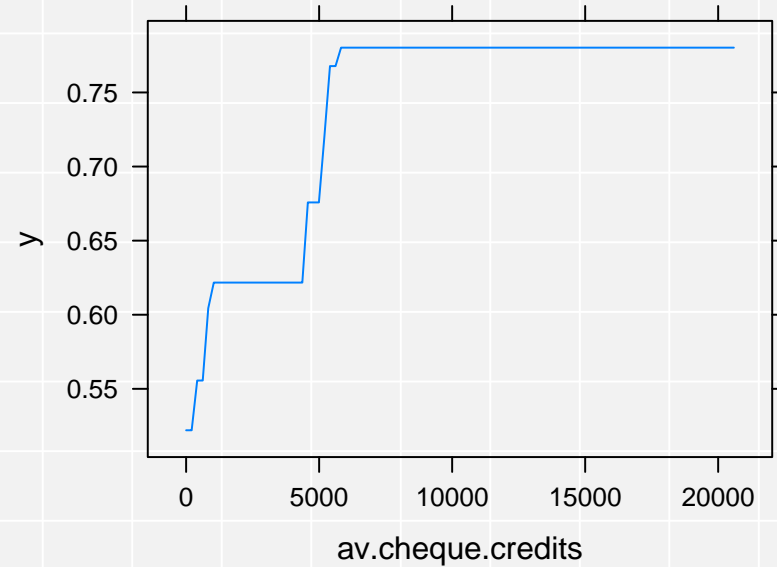
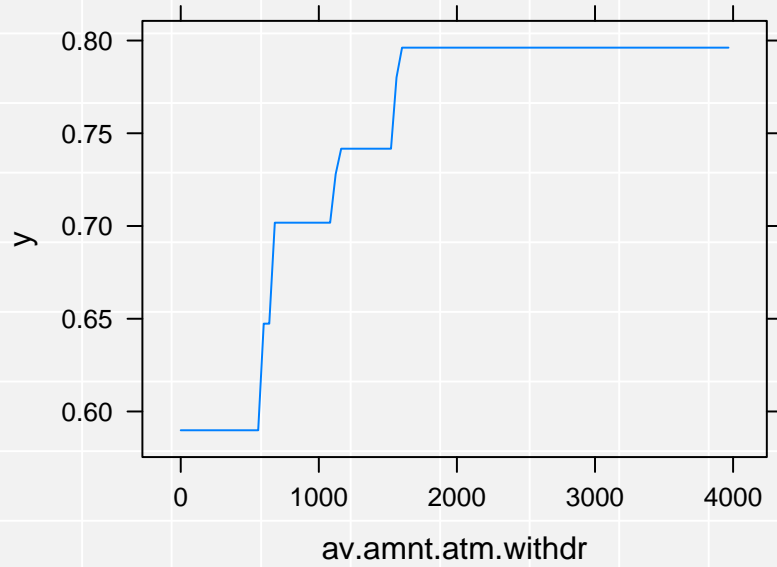
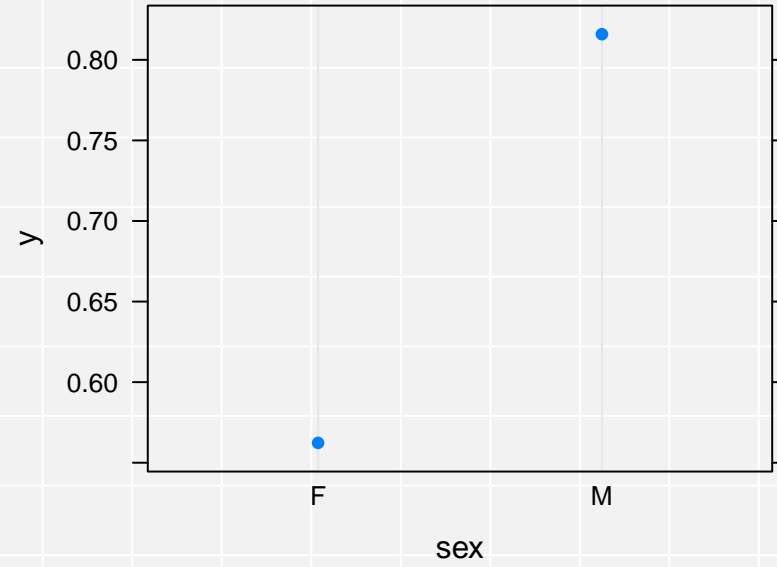

```
par(mar = c(4,12,1,1), cex = 0.7)
tmp <- summary(CCgbm) ## tmp is now the data frame displayed above.
```



Partial dependency plots are also available.

Look at the four “most influential” variables

```
suppressPackageStartupMessages({  
  library(gridExtra)  
})  
(vars <- tmp$var[1:4])                                ## four "most influential" variables  
  
[1] "sd.amnt.atm.withdr" "sex"                            "av.amnt.atm.withdr"  
[4] "av.cheque.credits"  
  
plots <- lapply(vars, . %>% plot(CCgbm, i.var = ., type = "response"))  
grobs <- arrangeGrob(grobs = plots, layout_matrix = rbind(1:2, 3:4))  
plot(grobs)
```



4 The final reckoning

Now to see how things worked out this time. First a helper function

```
Class <- function(object, newdata, ...)  
  UseMethod("Class")  
  
Class.rpart <- function(object, newdata, ...)  
  predict(object, newdata, type = "class")  
  
Class.bagRpart <- function(object, newdata, ...)  
  predict(object, newdata)  
  
Class.randomForest <- Class.C5.0 <- predict  
  
Class.glm <- Class.mboost <- Class.gbm <- function(object, newdata, ...) {  
  ## only applies for binomial GLMs with symmetric links  
  suppressMessages(predict(object, newdata) > 0)  
}
```

The helper function *Class* streamlines things a bit:


```
errorRate <- function(tab) 100*(1 - sum(diag(tab))/sum(tab))
true <- CCTest$credit.card.owner  # $
(res <- sort(sapply(list(Tree = CCTree,           PrunedTree = CCPTree,
                    SimpleBagging = CCSBag,       BayesBagging = CCBBag,
                    SimpleBaggingP = CCSBagP,     BayesBaggingP = CCBBagP,
                    RandomForest = CCRf,          C5.0 = CCc50,
                    GradientBoosting = CCgbm,    BoostedGlm = CCglmboost,
                    NaiveGLM = CCGlmNaive,       BoostedTree = CCblackboost,
                    Glm_AIC = CCGlmAIC,          GlmGIC = CCGlmGIC,
                    Glm_BIC = CCGlmBIC),
                    function(x) errorRate(table(Class(x, CCTest), true))))))
```

GradientBoosting	RandomForest	SimpleBaggingP	BayesBagging
11.35802	11.85185	12.59259	13.20988
SimpleBagging	BayesBaggingP	BoostedTree	C5.0
13.33333	13.58025	14.56790	14.93827
Tree	Glm_AIC	NaiveGLM	Glm_BIC
16.04938	16.41975	16.91358	17.03704
PrunedTree	GlmGIC	BoostedGlm	
17.16049	17.16049	17.65432	


```

par(mar = c(3,8,3,1))
barplot(rev(res), horiz=TRUE, las=1, fill = pal_green2brown)
axis(3)

```



References

- Buehlmann, P. and T. Hothorn (2007). Boosting algorithms: Regularization, prediction and model fitting (with discussion). *Statistical Science* 22, 477–505.
- Greenwell, B., B. Boehmke, J. Cunningham, and G. Developers (2020). *gbm: Generalized Boosted Regression Models*. R package version 2.1.8.
- Kuhn, M. and R. Quinlan (2020). *C50: C5.0 Decision Trees and Rule-Based Models*. R package version 0.1.3.1.
- Liaw, A. and M. Wiener (2002). Classification and regression by `randomForest`. *R News* 2(3), 18–22.
- Quinlan, R. (1993). *C4.5: Programs for Machine Learning*. Morgan Kaufmann Publishers.
- Ripley., B. (2012). *tree: Classification and Regression Trees*. CRAN. R package version 1.0–29.

Session information

Date: 2021-01-29

- R version 4.0.3 (2020-10-10), x86_64-pc-linux-gnu
- Running under: Ubuntu 20.04.1 LTS
- Matrix products: default
- BLAS: /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.9.0
- LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.9.0
- Base packages: base, datasets, graphics, grDevices, methods, parallel, stats, utils
- Other packages: C50 0.1.3.1, doParallel 1.0.16, dplyr 1.0.3, english 1.2-5, forcats 0.5.1, foreach 1.5.1, gbm 2.1.8, GGally 2.1.0, ggplot2 3.3.3, ggthemes 4.2.4, gridExtra 2.3, haven 2.3.1, iterators 1.0.13, knitr 1.31, lattice 0.20-41, lme4 1.1-26, Matrix 1.3-2, mboost 2.9-4, mgcv 1.8-33, microbenchmark 1.4-7, nlme 3.1-151, patchwork 1.1.1, purrr 0.3.4, randomForest 4.6-14, rbenchmark 1.0.0, Rcpp 1.0.6, readr 1.4.0, rpart 4.1-15, scales 1.1.1, SOAR 0.99-11, stabs 0.6-3, stringr 1.4.0, tibble 3.0.5, tidyr 1.1.2, tidyverse 1.3.0, visreg 2.7.0, WWRCourse 0.2.3, WWRData 0.1.0, WWRGraphics 0.1.2, WWRUtilities 0.1.2, xtable 1.8-4
- Loaded via a namespace (and not attached): assertthat 0.2.1, backports 1.2.1, boot 1.3-26, broom 0.7.3, cellranger 1.1.0, cli 2.2.0, codetools 0.2-18, colorspace 2.0-0, compiler 4.0.3, crayon 1.3.4, Cubist 0.2.3, DBI 1.1.1, dbplyr 2.0.0, digest 0.6.27, ellipsis 0.3.1, evaluate 0.14, fansi 0.4.2, farver 2.0.3, Formula 1.2-4, fractional 0.1.3, fs 1.5.0, generics 0.1.0, glue 1.4.2, grid 4.0.3, gtable 0.3.0, highr 0.8, hms 1.0.0, httr 1.4.2, inum 1.0-1, jsonlite 1.7.2, labeling 0.4.2, lazyData 1.1.0, libcoin 1.0-7, lifecycle 0.2.0, lubridate 1.7.9.2, magrittr 2.0.1, MASS 7.3-53, minqa 1.2.4, modelr 0.1.8, munsell 0.5.0, mvtnorm 1.1-1, nloptr 1.2.2.2, nnls 1.4, partykit 1.2-11, PBSmapping 2.73.0, pillar 1.4.7, pkgconfig 2.0.3, plyr 1.8.6, quadprog 1.5-8, R6 2.5.0, RColorBrewer 1.1-2, readxl 1.3.1, reprex 1.0.0, reshape 0.8.8, reshape2 1.4.4, rlang 0.4.10, rstudioapi 0.13, rvest 0.3.6, splines 4.0.3, statmod 1.4.35, stringi 1.5.3, survival 3.2-7, tidyselect 1.1.0, tools 4.0.3, vctrs 0.3.6, withr 2.4.1, xfun 0.20, xml2 1.3.2