

Working with 

A University of Queensland Advanced Workshop

Session 10: Some Programming Techniques

Bill Venables, CSIRO/Data61, Dutton Park

Rhetta Chappell, Griffith University

2–5 February, 2021

Contents

1	A simulation example	2
1.1	A simulation strategy	4
1.2	S3 methods	7
1.3	Some lessons	11
2	Intermission: On text processing	12
2.1	The Authorised Version of King James	12
2.2	Reading in	13
2.3	Phrase length	18
3	Convolutions	23
	Session information	38

1 A simulation example

Find by simulation the relative volume of a sphere of diameter 2 to a square box of side 2?

For dimensions $d = 1, 2, 3, \dots$

The answer is known mathematically. The volume of the sphere of radius r is

$$V_d^{(S)} = \frac{\pi^{d/2} r^d}{\Gamma(d/2 + 1)}$$

and the volume of the square box of side $2r$ is clearly $V_c^{(B)} = 2^d r^d$. Hence the relative size is

$$R =_{\text{def.}} \frac{V_d^{(S)}}{V_d^{(B)}} = \frac{\pi^{d/2}}{\Gamma(d/2 + 1) 2^d}$$

```

rvball <- local({
  pi <- base::pi  ## make sure it is the right one

  function(d)
    structure(exp(d/2 * log(pi) - d * log(2) - lgamma(d/2 + 1)),
              names = as.character(d))
})
floor(1/rvball(1:10))

```

1	2	3	4	5	6	7	8	9	10
1	1	1	3	6	12	27	63	155	401

1.1 A simulation strategy

- Generate samples uniformly within the box centred at $(0, 0, \dots, 0)$
- Count the number at a Euclidean distance no further than 1 from the origin.
- Take ratio.

```

mcrvball <- function(d, N = 100000, blocksize = 10000) {
  n2 <- inside <- 0
  while(n2 < N) {
    n1 <- n2
    n2 <- min(n2 + blocksize, N)
    No <- n2 - n1
    samp <- matrix(runif(No * d, -1, 1), No, d)
    inside <- inside + sum(rowSums(samp^2) < 1)
  }
  res <- list(dimensions = d, inside = inside,
             total = N, call = match.call())
  class(res) <- "mcrvball"
  res
}

```

An alternative implementation using the `doParallel` family of packages.

```
mcrvball2 <- function(d, N = 100000, blocksize = 10000) {  
  chunks <- idiv(N, chunkSize = blocksize) ## division iterator  
  inside <- foreach(No = chunks, .combine = sum) %dopar% {  
    samp <- matrix(runif(No*d, -1, 1), No, d)  
    sum(rowSums(samp^2) < 1)  
  }  
  structure(list(dimensions = d, inside = inside,  
                total = N, call = match.call()),  
            class = "mcrvball")  
}
```

To use it:

```
suppressPackageStartupMessages(library(doParallel))  
cl <- makePSOCKcluster(detectCores() - 1)  
registerDoParallel(cl) ## not needed in "parallel" package  
  
tst <- mcrvball2(10, N = 1e7, blocksize = 1e5)  
  
stopCluster(cl)  
rm(cl)
```

1.2 S3 methods

```
print.mcrvball <- function(x, ...) {  
  with(x, cat("Dim.:", dimensions,  
             "Estimated:", signif(inside/total, 4),  
             "Actual:", signif(rvball(dimensions), 4), "\n"))  
  invisible(x)  
}  
  
Ops.mcrvball <- function(e1, e2) { ## group generic  
  if(inherits(e1, "mcrvball"))  
    e1 <- with(e1, inside/total)  
  
  if(!missing(e2) && inherits(e2, "mcrvball"))  
    e2 <- with(e2, inside/total)  
  NextMethod()  
}
```



```
for(i in 4:10) print(mcrvball(i, 1000000))
```

```
Dim.: 4 Estimated: 0.3087 Actual: 0.3084
```

```
Dim.: 5 Estimated: 0.1645 Actual: 0.1645
```

```
Dim.: 6 Estimated: 0.08091 Actual: 0.08075
```

```
Dim.: 7 Estimated: 0.03692 Actual: 0.03691
```

```
Dim.: 8 Estimated: 0.01587 Actual: 0.01585
```

```
Dim.: 9 Estimated: 0.006474 Actual: 0.006442
```

```
Dim.: 10 Estimated: 0.002508 Actual: 0.00249
```

```
r <- numeric(7)
```

```
for(d in 4:10) r[d-3] <- floor(1/mcrvball(d)) ; r
```

```
[1] 3 6 12 27 62 160 390
```

Having a *call* component allows *update* to be used:

```
p10 <- mcrvball(10)
p10a <- update(p10, N = 1000000, blocksize = 100000)
c(1/p10, 1/p10a)

[1] 411.5226 402.7386

"%+%" <- function(e1, e2)
  UseMethod("%+%", e1)
"%+%.mcrvball" <- function(e1, e2) {
  if(e1$dimensions != e2$dimensions) stop("ball dimensions differ!")
  res <- list(dimensions = e1$dimensions, inside = e1$inside + e2$inside,
             total = e1$total + e2$total, call = e1$call)
  class(res) <- "mcrvball"
  res
}
# p10 %+% p10a
floor(1/(p10 %+% p10a))

[1] 403
```

Automatic (lazy) vectorization

```
Mcrvball <- Vectorize(mcrvball, vectorize.args = c("d", "N"), SIMPLIFY = FALSE)
Mcrvball(4:8, N = 1e6, blocksize = 1e5)
```

```
[[1]]
```

```
Dim.: 4 Estimated: 0.3082 Actual: 0.3084
```

```
[[2]]
```

```
Dim.: 5 Estimated: 0.1645 Actual: 0.1645
```

```
[[3]]
```

```
Dim.: 6 Estimated: 0.08064 Actual: 0.08075
```

```
[[4]]
```

```
Dim.: 7 Estimated: 0.03692 Actual: 0.03691
```

```
[[5]]
```

```
Dim.: 8 Estimated: 0.01581 Actual: 0.01585
```

1.3 Some lessons

- Vectorization. (*rvball*) and vectorization tools: (*Vectorize*)
- Taking the “whole object view” of the problem. (*mcrvball*)
- Object orientation: put all the information you are likely to need into the object and give it a class. (*mcrvball*)
- Methods for existing generics. (*print.mcrvball*)
- Group generic functions. (*Ops.mcrvball*)
- Binary operators and new generic functions. (%+%)

2 Intermission: On text processing

Data in the form of text is becoming more common and important.

R has many sophisticated tools and packages for text manipulation—here we just glimpse two elementary ones: *grep* and *(g)sub*: the first for finding, the second twin pair for fixing.

2.1 The Authorised Version of King James

The `extdata` subdirectory of the `WWRCourse` package folder contains two compressed text files, of the first and last books of AV:

```
subdir <- system.file("extdata", package = "WWRCourse")
dir(subdir, pattern = "\\\\.txt.gz$")      ## NB regular expression

[1] "Genesis.txt.gz"      "Revelations.txt.gz"
```

Our task is to compare *chapter length* *phrase length* and *word length* in these two document, as an aspect of literary style

2.2 Reading in

```
gen <- file.path(subdir, "Genesis.txt.gz")
rev <- file.path(subdir, "Revelations.txt.gz")
Genesis <- scan(gzfile(gen), what = "") ## no need to gunzip()
Revelations <- scan(gzfile(rev), what = "")
rbind(head(Genesis), head(Revelations)) %>% noquote

      [,1] [,2] [,3]      [,4]      [,5] [,6]
[1,] Gen.1 In   the      beginning God   created
[2,] Rev.1 The  Revelation of      Jesus Christ,
```

Both have chapter markers in them of the form Gen.1, Rev.23, which are not part of the text. We can get the chapter lengths from them, though.

```
grep("^Gen\\.\\.\\.[:digit:]]+$", Genesis)[1:5] ## the idea

[1]      1   799 1432 2128 2761

chapterLengths <- function(txt, ## cement it in a function & check
                             book = substring(deparse(substitute(txt)), 1, 3)) {
  regex <- paste0("^", book, "\\\\.[:digit:]]+$")
  where <- grep(regex, txt)
  diff(where - seq_along(where))
}
```

```

}
GenC <- chapterLengths(Genesis)
RevC <- chapterLengths(Revelations)
rbind(Gen = c(n = length(GenC), median = median(GenC)),
      Rev = c(n = length(RevC), median = median(RevC)))

```


	n	median
Gen	49	695
Rev	21	540

Revelations has fewer, and shorter chapters, but the chapter lengths are quite similar to the opening chapter lengths of Genesis:

```


plot(GenC, type="h", col = "#418A78", xlab = "Chapter #",
     ylab="Words", ylim = range(0, GenC, RevC))
points(RevC, pch=20, col="#FC7115", cex=1.5)
legend("topright", c("Genesis", "Revelations"), lty = c("solid", NA),
     pch=c(NA, 20), col = c("#418A78", "#FC7115"), bty="n")

```



A (contrived) *ggplot* version

```
gendat <- data.frame(x = c(1, rep(seq_along(GenC), each = 3)),  
                    y = c(0, rbind(0, GenC, 0)))  
revdat <- data.frame(x = seq_along(RevC), y = RevC)  
ggplot() + aes(x,y) + geom_path(data = gendat, colour = "#418A78") +  
  geom_point(data = revdat, colour = "#FC7115", size = 2) +  
  geom_hline(yintercept = 0, colour = "grey", size = 0.5) +  
  labs(x = "Chapter #", y = "Words") +  
  annotate("text", x = 50, y = 1800, label = "Genesis",  
          colour = "#418A78", hjust = 1) +  
  annotate("text", x = 50, y = 1750, label = "Revelations",  
          colour = "#FC7115", hjust = 1) + theme_bw()
```



2.3 Phrase length

The first step is to strip out the chapter markers and leave the words, and punctuation, only.

```
Gen <- grep("^Gen\\.\\.[:digit:]+$", Genesis, invert = TRUE, value = TRUE)
Rev <- grep("^Rev\\.\\.[:digit:]+$", Revelations, invert = TRUE, value = TRUE)
noquote(rbind(head(Gen), head(Rev)))
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	In	the	beginning	God	created	the
[2,]	The	Revelation	of	Jesus	Christ,	which

We define a phrase as the words between successive (terminal) punctuation marks. These indicate a pause in the reading.

```
GenP <- diff(c(0, grep("[:punct:]+$", Gen)))
RevP <- diff(c(0, grep("[:punct:]+$", Rev)))
rbind(Gen = head(GenP), Rev = head(RevP))
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
Gen	10	6	2	9	12	3
Rev	5	5	12	13	8	7

Put into a data frame and look at the relative frequencies.^a


```
Phrase <- rbind(data.frame(Book = "Gen", Length = GenP),
                data.frame(Book = "Rev", Length = RevP))
tab <- with(Phrase, table(Book, Length)); t(tab)
```

	Book	
Length	Gen	Rev
1	400	108
2	537	167
...		
23	0	1
24	0	1
25	1	0
26	0	1

```
colours <- c(Genesis = "#418A78", Revelations = "#FC7115")
par(cex.axis = 0.8)
barplot(tab/rowSums(tab), beside = TRUE, main = "Phrase Length",
        fill = colours, colour = colours)
grid()
```

^aExercise: Assuming Length-1 is Poisson, test for a difference in mean phrase length between the two books

Phrase Length



Finally we come to word length distributions. For this we strip out any non-letter character and count the string lengths of what is left.

```
GenW <- nchar(gsub("[^[:alpha:]]", "", Gen))
RevW <- nchar(gsub("[^[:alpha:]]", "", Rev))
noquote(rbind(head(Rev), head(RevW))) ## word 5 comma excluded
```


	[,1]	[,2]		[,3]	[,4]	[,5]	[,6]
[1,]	The	Revelation	of	Jesus	Christ,	which	
[2,]	3	10	2	5	6	5	

```
Words <- rbind(data.frame(Book = "Gen", Length = GenW),
               data.frame(Book = "Rev", Length = RevW))
tab <- with(Words, table(Book, Length)); t(tab)
```

	Book	
Length	Gen	Rev
....		
14	7	2
15	2	1

```
par(cex.axis = 0.8)
barplot(tab/rowSums(tab), beside = TRUE, main = "Word Length",
        fill = colours, colour = colours)
grid()
```

Word Length



3 Convolutions

Given two sequences of numbers, $a_i, i = 0, 1, \dots$ and $b_j, j = 0, 1, \dots$ their *convolution* is defined by

$$(ab)_k = \sum_{i+j=k} a_i b_j$$

(this is the operation involved in polynomial multiplication).

Consider some methods for doing this in **R**.


```

convolve0 <- function(a, b) {
  ab <- rep(0, length(a) + length(b) - 1)
  for(i in 1:length(a))
    for(j in 1:length(b))
      ab[i+j-1] <- ab[i+j-1] + a[i]*b[j]
  ab
}
###
convolve1 <- function(a, b) {
  ab <- rep(0, length(a) + length(b) - 1)
  ind <- 1:length(a)
  for(j in 1:length(b)) {
    ab[ind] <- ab[ind] + a*b[j]
    ind <- ind + 1
  }
  ab
}

```

```

convolve1a <- function(a, b) {
  if(length(a) < length(b)) Recall(b,a) else {
    ab <- rep(0, length(a) + length(b) - 1)
    ind <- 1:length(a)
    for(j in 1:length(b)) {
      ab[ind] <- ab[ind] + a*b[j]
      ind <- ind + 1
    }
    ab
  }
}

###
convolve2 <- function(a, b) {
  p <- outer(a, b)
  as.vector(tapply(p, row(p) + col(p), sum))
}

### Uses much less memory than convolve2
convolve2a <- function(a, b)
  as.vector(tapply(outer(a, b),
    outer(seq(along = a), seq(along = b), "+"), sum))

```

The young geek's version.

```
convolve_hw <- function(a, b) {  
  stopifnot(require(dplyr))  
  data.frame(x = as.vector(outer(a, b)),  
             g = as.vector(outer(seq_along(a),  
                                 seq_along(b), "+"))) %>%  
    group_by(g) %>%  
    summarise(conv = sum(x), .groups = 'drop') %>%  
    .[["conv"]]  
}
```

A C code version

File VR_convolve.c has the code:

```
void VR_convolve(double *a, int *na,
                 double *b, int *nb,
                 double *ab)
{
    int i, j, nab = *na + *nb - 1;

    for(i = 0; i < nab; i++) ab[i] = 0.0;
    for(i = 0; i < *na; i++)
        for(j = 0; j < *nb; j++)
            ab[i + j] += a[i] * b[j];
}
```

To compile:

```
$ R CMD SHLIB VR_convolve.c
```

```

convolve3 <- function(a, b) {
  if(!is.loaded("VR_convolve")) {
    path <- file.path("SharedObjects",
                      paste("VR_convolve",
                            .Platform$dynlib.ext, sep=""))

    dyn.load(path)
  }
  storage.mode(a) <- "double"
  storage.mode(b) <- "double"

  .C("VR_convolve",
     a,
     length(a),
     b, length(b),
     ab = double(length(a) + length(b) - 1))$ab
}

```

The Rcpp revolution

- Makes it easy to write compiled code *without* the need for an R-side interface function (such as the above).
- Can interact directly with R-objects, using R to make *most* of the mode conversions
- Good tools for creating packages using such code, (*Rcpp.package.skeleton*, in particular).
- Drawbacks:
 - The code has to be written specifically for Rcpp, using the extensive C++ header files,
 - There *can be* a small performance overhead using “RcppSugar”, (but quicker programming and fewer bugs).

Use Rcpp to extend the R system; Use the *dyn.load* to use of existing code, or minimally modified.

An `Rcpp` convolution function: file `src/convolve3a.Cpp`

```
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]
NumericVector convolve3a(NumericVector x, NumericVector y)
{
    int nx = x.size(), ny = y.size(), nz = nx + ny - 1;
    NumericVector z(nz); // set to 0 on creation. NB z() not z[] here!

    for(int i = 0; i < nx; ++i) {
        for(int j = 0; j < ny; ++j) {
            z[i+j] += x[i]*y[j];
        }
    }
    return z;
}
```

To make the code available as an **R** function, you need to have the necessary tools installed, but then:

```
library(Rcpp)
sourceCpp("src/convolve3a.cpp")
```

The **R** version of the function looks like:

```
convolve3a

function (x, y)
.Call(<pointer: 0x7f60987ca630>, x, y)
```


Alternatively, small functions can be compiled directly as a text string. The headers are assumed:

```
Rcpp::cppFunction('
NumericVector convolve3a(NumericVector x, NumericVector y)
{
    int nx = x.size(), ny = y.size(), nz = nx + ny - 1;
    NumericVector z(nz); // set to 0 on creation. NB z() not z[] here!
    for(int i = 0; i < nx; ++i) {
        for(int j = 0; j < ny; ++j) {
            z[i+j] += x[i]*y[j];
        }
    }
    return z;
}
')
```

Such functions *must* be re-made the first time they are needed in an **R** session, in order to integrate them with the **R** program. They cannot be *save()*d in one session and *load()*ed in another, for example.

Permanent versions *can* be made in packages, however, with the system integration automatically happening when the package is loaded.

Some checks:

```
a <- 1:3; b <- 4:7  
rbind(convolve0(a,b), convolve1(a,b), convolve1a(a,b),  
      convolve2(a,b), convolve2a(a,b), convolve3(a,b),  
      convolve3a(a,b), convolve_hw(a, b))
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	4	13	28	34	32	21
[2,]	4	13	28	34	32	21
[3,]	4	13	28	34	32	21
[4,]	4	13	28	34	32	21
[5,]	4	13	28	34	32	21
[6,]	4	13	28	34	32	21
[7,]	4	13	28	34	32	21
[8,]	4	13	28	34	32	21

```

library(microbenchmark)
a <- 1:300; b <- 4:7
(b <- microbenchmark(convolve_hw(a, b),
                      convolve0(a,b),
                      convolve1(a,b),
                      convolve1a(a,b),
                      convolve2(a,b),
                      convolve2a(a,b),
                      convolve3(a,b),
                      convolve3a(a,b))) %>% summary() %>% arrange(median) %>%
  .[, cs(expr, min, median, mean, max, cld)]

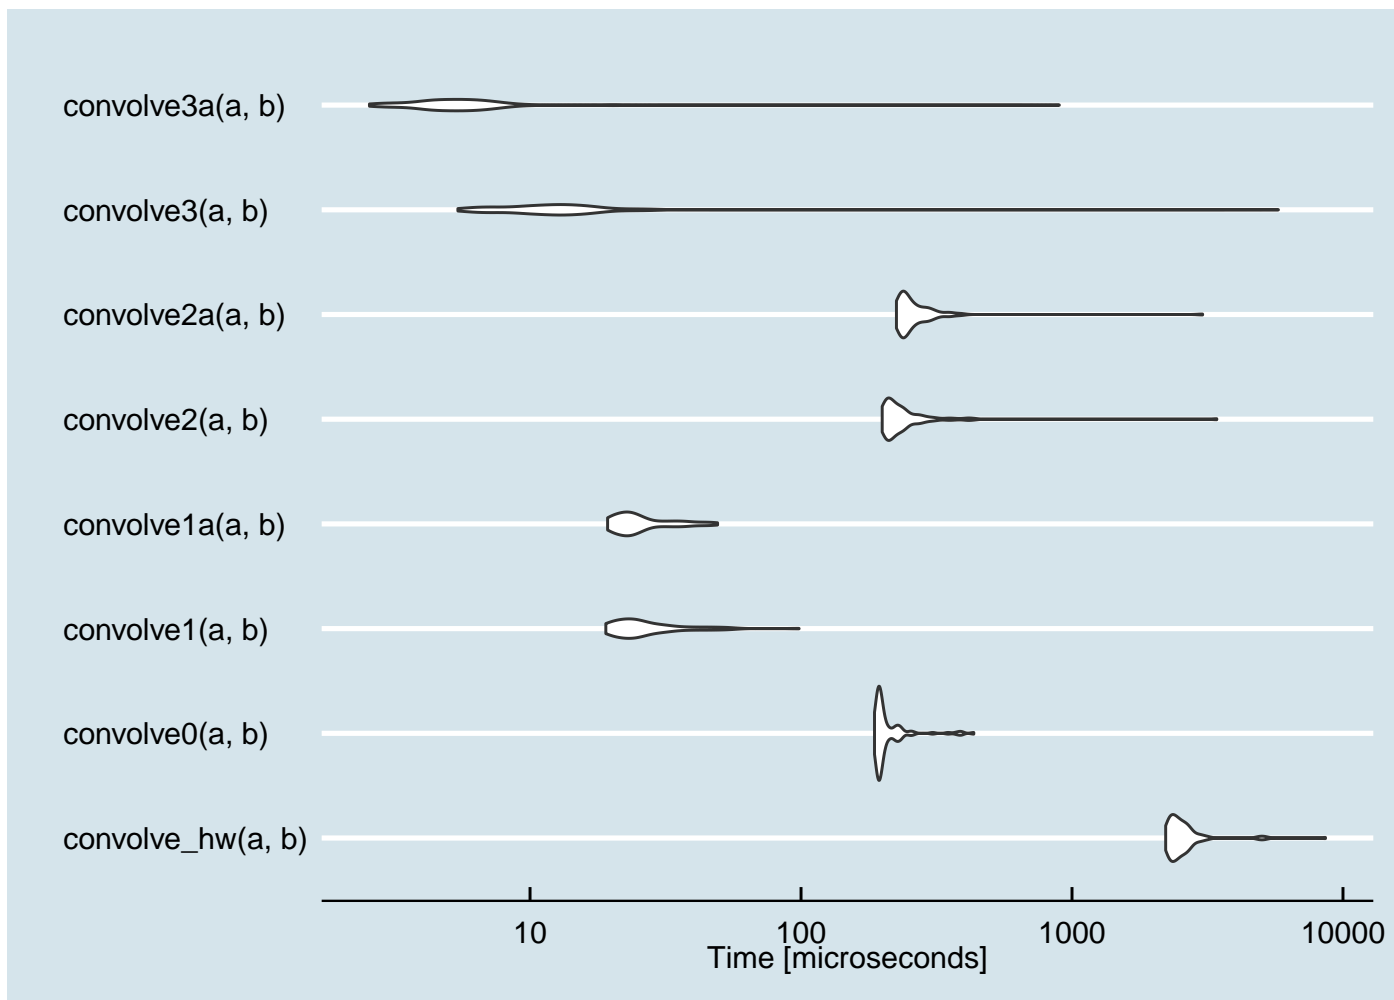
```

	expr	min	median	mean	max	cld
1	convolve3a(a, b)	2.555	5.3060	14.72706	897.728	a
2	convolve3(a, b)	5.417	12.1750	70.08196	5776.490	ab
3	convolve1a(a, b)	19.331	23.7105	26.52728	49.217	a
4	convolve1(a, b)	19.043	24.2050	28.04480	98.359	a
5	convolve0(a, b)	186.723	195.9835	212.63153	434.143	bc
6	convolve2(a, b)	199.330	221.9620	270.35601	3422.914	c
7	convolve2a(a, b)	224.943	246.1995	288.69014	3034.108	c
8	convolve_hw(a, b)	2216.659	2459.1560	2660.45093	8618.315	d

```

suppressMessages(autoplot(b)) + theme_economist()

```



Some timings (the old fashioned way):

```
a <- 1:1000; b <- 1:10000
library(rbenchmark)
benchmark(convolve0 (a,b), convolve1 (a,b), convolve1 (b,a),
          convolve1a(a,b), convolve1a(b,a), convolve2 (a,b),
          convolve2a(a,b), convolve3 (a,b), convolve3a(a,b),
          columns = c("test", "replications", "elapsed", "relative"),
          order = "relative", relative = "elapsed", replications = 20)
```

	test	replications	elapsed	relative
9	convolve3a(a, b)	20	0.096	1.000
8	convolve3(a, b)	20	0.099	1.031
3	convolve1(b, a)	20	2.804	29.208
4	convolve1a(a, b)	20	2.810	29.271
5	convolve1a(b, a)	20	2.948	30.708
2	convolve1(a, b)	20	3.068	31.958
6	convolve2(a, b)	20	16.996	177.042
7	convolve2a(a, b)	20	17.492	182.208
1	convolve0(a, b)	20	22.478	234.146

Session information

Date: 2021-01-29

- R version 4.0.3 (2020-10-10), x86_64-pc-linux-gnu
- Running under: Ubuntu 20.04.1 LTS
- Matrix products: default
- BLAS: /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.9.0
- LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.9.0
- Base packages: base, datasets, graphics, grDevices, methods, parallel, stats, utils
- Other packages: doParallel 1.0.16, dplyr 1.0.3, english 1.2-5, forcats 0.5.1, foreach 1.5.1, ggplot2 3.3.3, ggthemes 4.2.4, gridExtra 2.3, iterators 1.0.13, knitr 1.31, lattice 0.20-41, microbenchmark 1.4-7, patchwork 1.1.1, purrr 0.3.4, Rcpp 1.0.6, readr 1.4.0, scales 1.1.1, stringr 1.4.0, tibble 3.0.5, tidyr 1.1.2, tidyverse 1.3.0, WWRCourse 0.2.3, WWRData 0.1.0, WWRGraphics 0.1.2, WWRUtilities 0.1.2, xtable 1.8-4
- Loaded via a namespace (and not attached): assertthat 0.2.1, backports 1.2.1, broom 0.7.3, cellranger 1.1.0, cli 2.2.0, codetools 0.2-18, colorspace 2.0-0, compiler 4.0.3, crayon 1.3.4, DBI 1.1.1, dbplyr 2.0.0, digest 0.6.27, ellipsis 0.3.1, evaluate 0.14, fansi 0.4.2, farver 2.0.3, fractional 0.1.3, fs 1.5.0, generics 0.1.0, glue 1.4.2, grid 4.0.3, gtable 0.3.0, haven 2.3.1, highr 0.8, hms 1.0.0, http 1.4.2, jsonlite 1.7.2, labeling 0.4.2, lazyData 1.1.0, lifecycle 0.2.0, lubridate 1.7.9.2, magrittr 2.0.1, MASS 7.3-53, Matrix 1.3-2, modelr 0.1.8, multcomp 1.4-15, munsell 0.5.0, mvtnorm 1.1-1, PBSmapping 2.73.0, pillar 1.4.7, pkgconfig 2.0.3, R6 2.5.0, randomForest 4.6-14, readxl 1.3.1, reprex 1.0.0, rlang 0.4.10, rpart 4.1-15, rstudioapi 0.13, rvest 0.3.6, sandwich 3.0-0, SOAR 0.99-11, splines 4.0.3, stringi 1.5.3, survival 3.2-7, TH.data 1.0-10, tidyselect 1.1.0, tools 4.0.3, vctrs 0.3.6, withr 2.4.1, xfun 0.20, xml2 1.3.2, zoo 1.8-8